

# Parallel Programming in C with MPI and OpenMP



# Parallel Programming in C with MPI and OpenMP

Michael J. Quinn





# Algorithm design and basic algorithms

Slides are modified from those found in  
*Parallel Programming in C with MPI and  
OpenMP*, Michael Quinn

# Outline

# Outline

- Task/channel model

# Outline

- Task/channel model
- Algorithm design methodology

# Outline

- Task/channel model
- Algorithm design methodology
- Case studies



# Task/Channel Model

# Task/Channel Model

- Parallel computation = set of tasks

# Task/Channel Model

- Parallel computation = set of tasks
- Task

# Task/Channel Model

- Parallel computation = set of tasks
- Task
  - ◆ Program

# Task/Channel Model

- Parallel computation = set of tasks
- Task
  - ◆ Program
  - ◆ Local memory


# Task/Channel Model

- Parallel computation = set of tasks
- Task
  - ◆ Program
  - ◆ Local memory
  - ◆ Collection of I/O ports

# Task/Channel Model

- Parallel computation = set of tasks
- Task
  - ◆ Program
  - ◆ Local memory
  - ◆ Collection of I/O ports
- Tasks interact by sending messages through channels

# Task/Channel Model





# Foster's Design Methodology

# Foster's Design Methodology

## ■ Partitioning

# Foster's Design Methodology

- Partitioning
- Communication

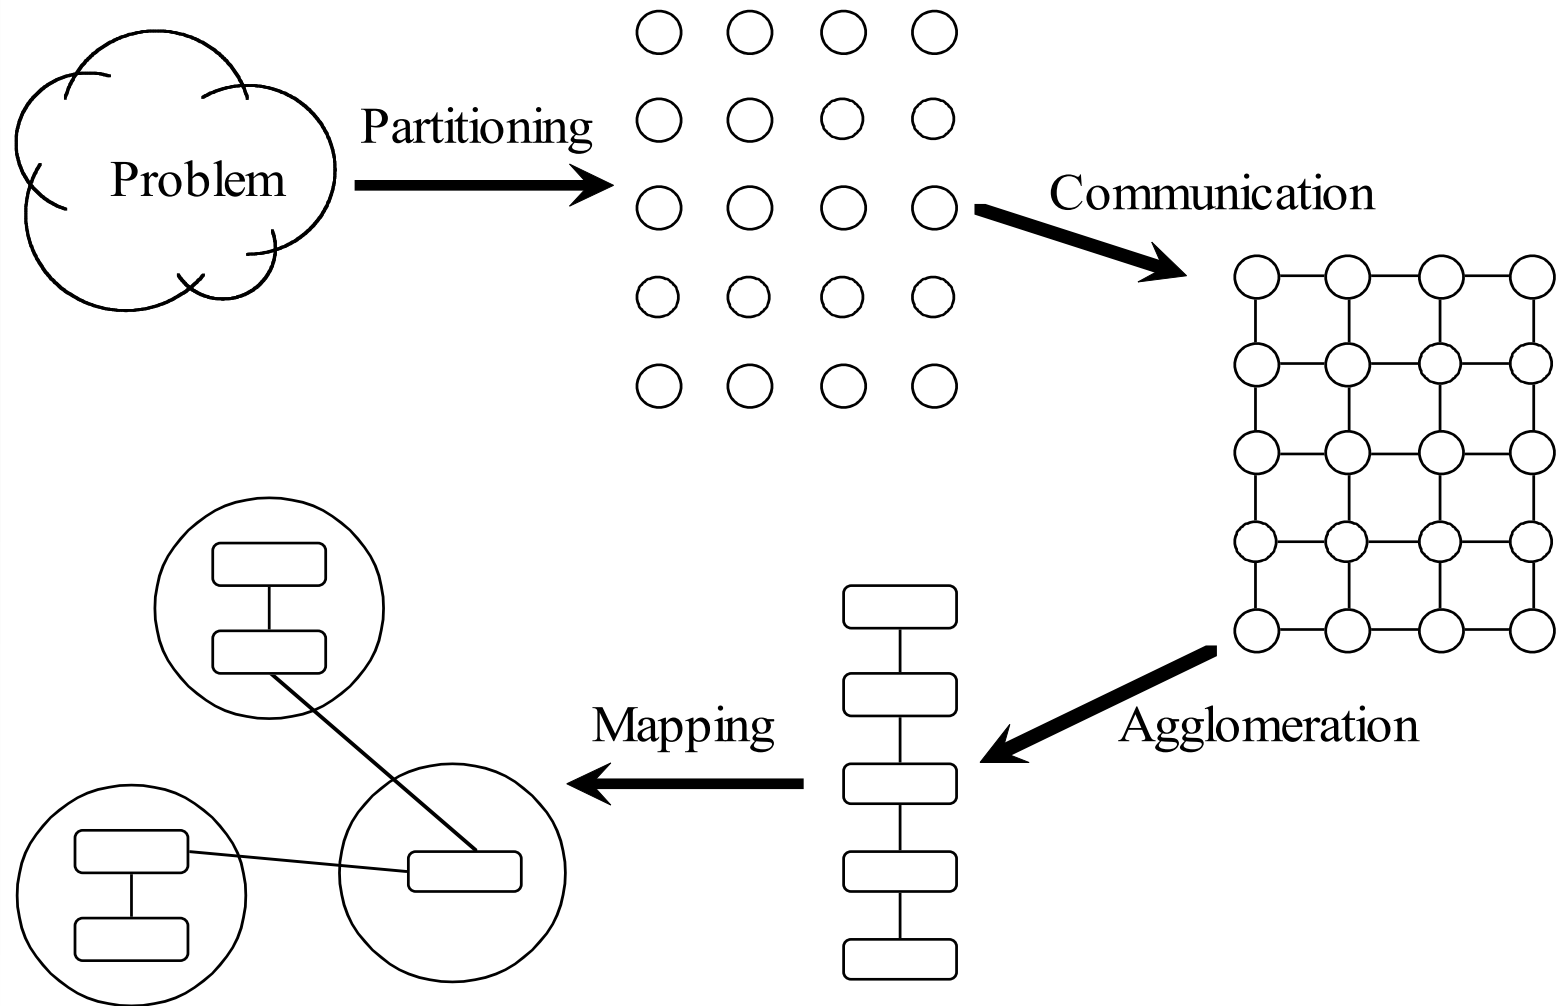
# Foster's Design Methodology

- Partitioning
- Communication
- Agglomeration

# Foster's Design Methodology

- Partitioning
- Communication
- Agglomeration
- Mapping

# Foster's Methodology



# Partitioning

# Partitioning

- Dividing computation and data into pieces



# Partitioning

- Dividing computation and data into pieces
- Domain decomposition

# Partitioning

- Dividing computation and data into pieces
- Domain decomposition
  - ◆ Divide data into pieces

# Partitioning

- Dividing computation and data into pieces
- Domain decomposition
  - ◆ Divide data into pieces
  - ◆ Determine how to associate computations with the data

# Partitioning

- Dividing computation and data into pieces
- Domain decomposition
  - ◆ Divide data into pieces
  - ◆ Determine how to associate computations with the data
- Functional decomposition


# Partitioning

- Dividing computation and data into pieces
- Domain decomposition
  - ◆ Divide data into pieces
  - ◆ Determine how to associate computations with the data
- Functional decomposition
  - ◆ Divide computation into pieces

# Partitioning

- Dividing computation and data into pieces
- Domain decomposition
  - ◆ Divide data into pieces
  - ◆ Determine how to associate computations with the data
- Functional decomposition
  - ◆ Divide computation into pieces
  - ◆ Determine how to associate data with the computations

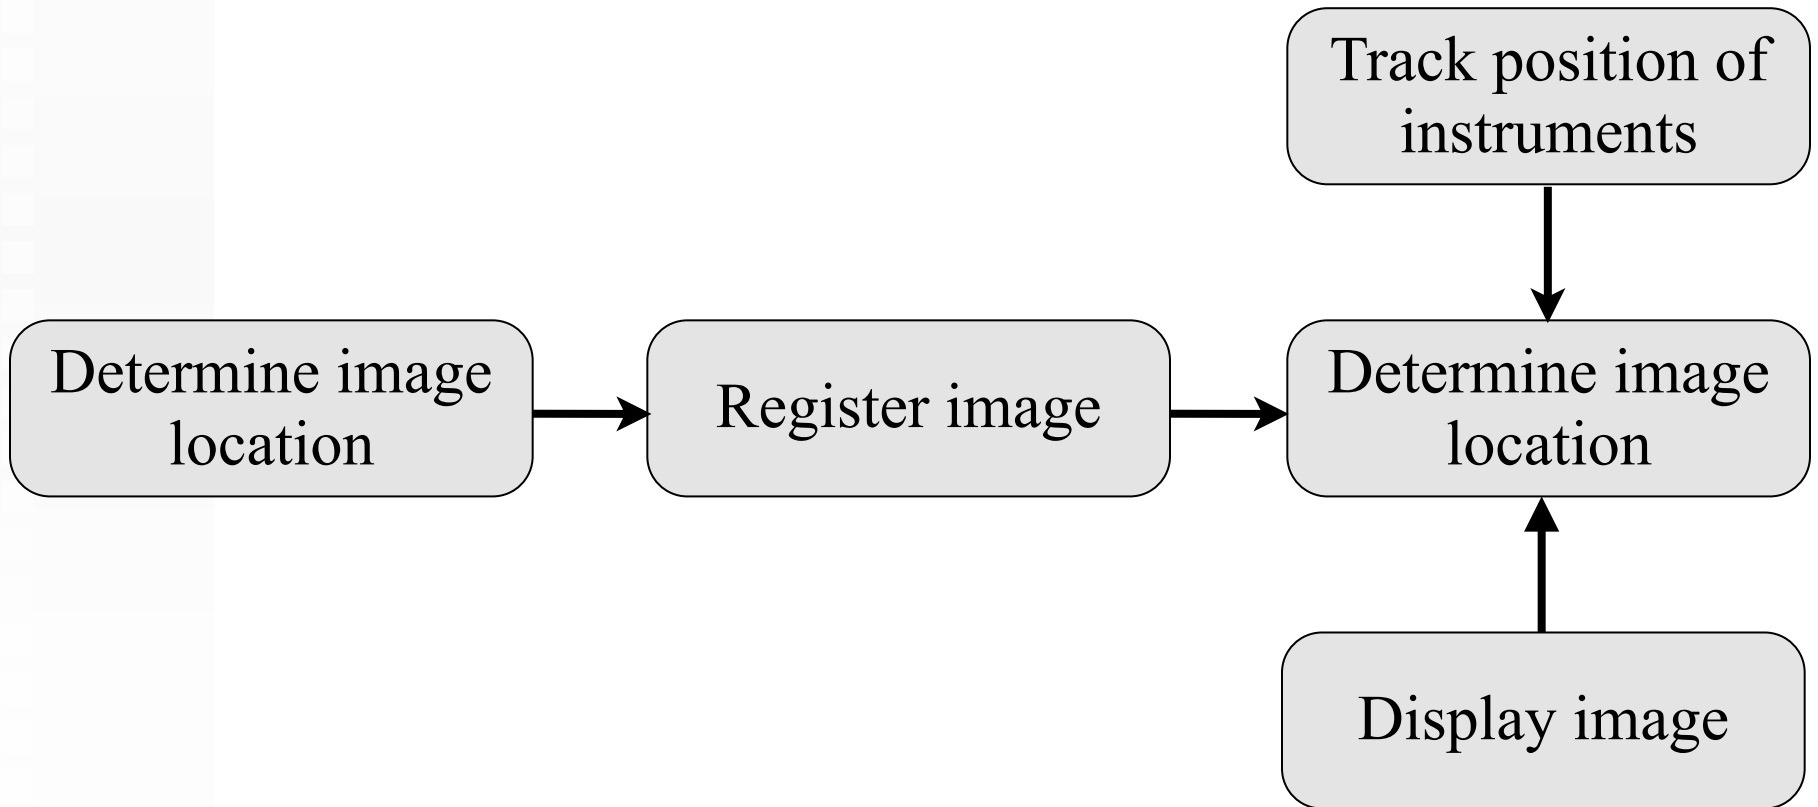
# Example Domain Decompositions



Primitive tasks is the number of scope, or order of magnitude, of the parallelism.

1-D has, in the example,  $n$ -way ||ism along the  $n$ -faces, 2-D has  $n^2$  ||ism along the faces, and 3-way has  $n^3$  ||ism along the faces.

# Example Functional Decomposition





# Types of parallelism

# Types of parallelism

- Numerical algorithms often have data-parallelism

# Types of parallelism

- Numerical algorithms often have data-parallelism
- Non-numerical algorithms often have functional parallelism.

# Types of parallelism

- Numerical algorithms often have data-parallelism
- Non-numerical algorithms often have functional parallelism.
- Many algorithms, especially complex numerical algorithms, have both, e.g., data parallelism within an function, many functions that can be done in parallel.

# Types of parallelism

- Numerical algorithms often have data-parallelism
- Non-numerical algorithms often have functional parallelism.
- Many algorithms, especially complex numerical algorithms, have both, e.g., data parallelism within an function, many functions that can be done in parallel.
- Functional parallelism often scales worse with increasing data size (concurrency-limited in isoefficiency terms)

# Partitioning Checklist

# Partitioning Checklist

- At least 10x more primitive tasks than processors in target computer

# Partitioning Checklist

- At least 10x more primitive tasks than processors in target computer
- Minimize redundant computations and redundant data storage



# Partitioning Checklist

- At least 10x more primitive tasks than processors in target computer
- Minimize redundant computations and redundant data storage
- Primitive tasks roughly the same size

# Partitioning Checklist

- At least 10x more primitive tasks than processors in target computer
- Minimize redundant computations and redundant data storage
- Primitive tasks roughly the same size
- Number of tasks an increasing function of problem size

# Communication

# Communication

- Determine values passed among tasks

# Communication

- Determine values passed among tasks
- Local communication

# Communication

- Determine values passed among tasks
- Local communication
  - ◆ Task needs values from a small number of other tasks

# Communication

- Determine values passed among tasks
- Local communication
  - ◆ Task needs values from a small number of other tasks
  - ◆ Create channels illustrating data flow

# Communication

- Determine values passed among tasks
- Local communication
  - ◆ Task needs values from a small number of other tasks
  - ◆ Create channels illustrating data flow
- Global communication



# Communication

- Determine values passed among tasks
- Local communication
  - ◆ Task needs values from a small number of other tasks
  - ◆ Create channels illustrating data flow
- Global communication
  - ◆ Significant number of tasks contribute data to perform a computation

# Communication

- Determine values passed among tasks
- Local communication
  - ◆ Task needs values from a small number of other tasks
  - ◆ Create channels illustrating data flow
- Global communication
  - ◆ Significant number of tasks contribute data to perform a computation
  - ◆ Don't create channels for them early in design

# Communication Checklist

# Communication Checklist

- Communication operations balanced among tasks

# Communication Checklist

- Communication operations balanced among tasks
- Each task communicates with only small group of neighbors

# Communication Checklist

- Communication operations balanced among tasks
- Each task communicates with only small group of neighbors
- Tasks can perform communications concurrently

# Communication Checklist

- Communication operations balanced among tasks
- Each task communicates with only small group of neighbors
- Tasks can perform communications concurrently
- Task can perform computations concurrently

# Agglomeration



# Agglomeration

- Grouping tasks into larger tasks

# Agglomeration

- Grouping tasks into larger tasks
- Goals

# Agglomeration

- Grouping tasks into larger tasks
- Goals
  - ◆ Improve performance

# Agglomeration

- Grouping tasks into larger tasks
- Goals
  - ◆ Improve performance
  - ◆ Maintain scalability of program


# Agglomeration

- Grouping tasks into larger tasks
- Goals
  - ◆ Improve performance
  - ◆ Maintain scalability of program
  - ◆ Simplify programming

# Agglomeration


- Grouping tasks into larger tasks
- Goals
  - ◆ Improve performance
  - ◆ Maintain scalability of program
  - ◆ Simplify programming
- In MPI programming, goal often to create one agglomerated task per processor

# Agglomeration Can Improve Performance



# Agglomeration Can Improve Performance


- Eliminate communication between primitive tasks agglomerated into consolidated task





# Agglomeration Can Improve Performance

- Eliminate communication between primitive tasks agglomerated into consolidated task
- Combine groups of sending and receiving tasks



# Agglomeration Checklist

# Agglomeration Checklist

- Locality of parallel algorithm has increased

# Agglomeration Checklist

- Locality of parallel algorithm has increased
- Replicated computations take less time than communications they replace

# Agglomeration Checklist

- Locality of parallel algorithm has increased
- Replicated computations take less time than communications they replace
- Data replication doesn't affect scalability

# Agglomeration Checklist

- Locality of parallel algorithm has increased
- Replicated computations take less time than communications they replace
- Data replication doesn't affect scalability
- Agglomerated tasks have similar computational and communications costs

# Agglomeration Checklist

- Locality of parallel algorithm has increased
- Replicated computations take less time than communications they replace
- Data replication doesn't affect scalability
- Agglomerated tasks have similar computational and communications costs
- Number of tasks increases with problem size

# Agglomeration Checklist

- Locality of parallel algorithm has increased
- Replicated computations take less time than communications they replace
- Data replication doesn't affect scalability
- Agglomerated tasks have similar computational and communications costs
- Number of tasks increases with problem size
- Number of tasks suitable for likely target systems



# Agglomeration Checklist

- Locality of parallel algorithm has increased
- Replicated computations take less time than communications they replace
- Data replication doesn't affect scalability
- Agglomerated tasks have similar computational and communications costs
- Number of tasks increases with problem size
- Number of tasks suitable for likely target systems
- Tradeoff between agglomeration and code modifications costs is reasonable

# Mapping

# Mapping

- Process of assigning tasks to processors

# Mapping

- Process of assigning tasks to processors
- Shared memory system: mapping done by operating system

# Mapping

- Process of assigning tasks to processors
- Shared memory system: mapping done by operating system
- Distributed memory system: mapping done by user

# Mapping

- Process of assigning tasks to processors
- Shared memory system: mapping done by operating system
- Distributed memory system: mapping done by user
- Conflicting goals of mapping

# Mapping


- Process of assigning tasks to processors
- Shared memory system: mapping done by operating system
- Distributed memory system: mapping done by user
- Conflicting goals of mapping
  - ◆ Maximize processor utilization

# Mapping

- Process of assigning tasks to processors
- Shared memory system: mapping done by operating system
- Distributed memory system: mapping done by user
- Conflicting goals of mapping
  - ◆ Maximize processor utilization
  - ◆ Minimize interprocessor communication



# Mapping Example



While this may reduce communication, load balance may be an issue

# Optimal Mapping

# Optimal Mapping

- Finding optimal mapping is NP-hard

# Optimal Mapping

- Finding optimal mapping is NP-hard
- Must rely on heuristics

# Optimal Mapping

- Finding optimal mapping is NP-hard
- Must rely on heuristics
- Metis is a popular package for partitioning graphs

# Optimal Mapping

- Finding optimal mapping is NP-hard
- Must rely on heuristics
- Metis is a popular package for partitioning graphs
  - ◆ Minimizes the number of edges between nodes in a graph

# Optimal Mapping

- Finding optimal mapping is NP-hard
- Must rely on heuristics
- Metis is a popular package for partitioning graphs
  - ◆ Minimizes the number of edges between nodes in a graph
  - ◆ Edges, for our purposes, can be thought of as communication

# Mapping Decision Tree



# Mapping Decision Tree

- Static number of tasks

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task
      - Agglomerate tasks to minimize communication

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task
      - Agglomerate tasks to minimize communication
      - Create one task per processor

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task
      - Agglomerate tasks to minimize communication
      - Create one task per processor
    - ◆ Variable computation time per task

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task
      - Agglomerate tasks to minimize communication
      - Create one task per processor
    - ◆ Variable computation time per task
      - Cyclically map tasks to processors

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task
      - Agglomerate tasks to minimize communication
      - Create one task per processor
    - ◆ Variable computation time per task
      - Cyclically map tasks to processors
      - GSS (guided self scheduling)



# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task
      - Agglomerate tasks to minimize communication
      - Create one task per processor
    - ◆ Variable computation time per task
      - Cyclically map tasks to processors
      - GSS (guided self scheduling)
  - ◆ Unstructured communication

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task
      - Agglomerate tasks to minimize communication
      - Create one task per processor
    - ◆ Variable computation time per task
      - Cyclically map tasks to processors
      - GSS (guided self scheduling)
  - ◆ Unstructured communication
    - Use a static load balancing algorithm

# Mapping Decision Tree

- Static number of tasks
  - ◆ Structured communication
    - ◆ Constant computation time per task
      - Agglomerate tasks to minimize communication
      - Create one task per processor
    - ◆ Variable computation time per task
      - Cyclically map tasks to processors
      - GSS (guided self scheduling)
  - ◆ Unstructured communication
    - Use a static load balancing algorithm
- Dynamic number of tasks

# Mapping Strategy

# Mapping Strategy

- Static number of tasks

# Mapping Strategy

- Static number of tasks
- Dynamic number of tasks

# Mapping Strategy

- Static number of tasks
- Dynamic number of tasks
  - ◆ Frequent communications between tasks

# Mapping Strategy

- Static number of tasks
- Dynamic number of tasks
  - ◆ Frequent communications between tasks
    - ◆ Use a dynamic load balancing algorithm



# Mapping Strategy

- Static number of tasks
- Dynamic number of tasks
  - ◆ Frequent communications between tasks
    - ◆ Use a dynamic load balancing algorithm
  - ◆ Many short-lived tasks

# Mapping Strategy

- Static number of tasks
- Dynamic number of tasks
  - ◆ Frequent communications between tasks
    - ◆ Use a dynamic load balancing algorithm
  - ◆ Many short-lived tasks
    - ◆ Use a run-time task-scheduling algorithm

# Mapping Strategy

- Static number of tasks
- Dynamic number of tasks
  - ◆ Frequent communications between tasks
    - ◆ Use a dynamic load balancing algorithm
  - ◆ Many short-lived tasks
    - ◆ Use a run-time task-scheduling algorithm
    - ◆ Cilk and Galois, discussed in the next couple of weeks, do this.

# Mapping Checklist

# Mapping Checklist

- Considered designs based on one task per processor and multiple tasks per processor

# Mapping Checklist

- Considered designs based on one task per processor and multiple tasks per processor
- Evaluated static and dynamic task allocation

# Mapping Checklist

- Considered designs based on one task per processor and multiple tasks per processor
- Evaluated static and dynamic task allocation
- If dynamic task allocation chosen, task allocator is not a bottleneck to performance

# Mapping Checklist

- Considered designs based on one task per processor and multiple tasks per processor
- Evaluated static and dynamic task allocation
- If dynamic task allocation chosen, task allocator is not a bottleneck to performance
- If static task allocation chosen, ratio of tasks to processors is at least 10:1



# Case Studies

# Case Studies

- Boundary value problem

# Case Studies

- Boundary value problem
- Finding the maximum


# Case Studies

- Boundary value problem
- Finding the maximum
- The n-body problem


# Case Studies

- Boundary value problem
- Finding the maximum
- The n-body problem
- Adding data input

# Boundary Value Problem





# Rod Cools as Time Progresses



# Want to use finite-difference method over multiple time steps


time




position



# Want to use finite-difference method over multiple time steps




# Want to use finite-difference method over multiple time steps



# Want to use finite-difference method over multiple time steps

time





Temperature at time  $t+1$  for a position on the rod represented by a node depends on the temperature of neighbors at time  $t$

position

# Partitioning

time





position



# Partitioning

- One data item per grid point

time





position




# Partitioning

time





- One data item per grid point
- Associate one primitive task with each grid point

position



# Partitioning

time





- One data item per grid point
- Associate one primitive task with each grid point
- Two-dimensional domain decomposition

position



# Communication


time



position








position

# Communication


time



- Identify communication pattern between primitive tasks
- Each interior primitive task has three incoming and three outgoing channels

position

# Agglomeration and Mapping



(a)




(b)



(c)


# Agglomeration and Mapping



(a)



(b)



(c)

Agglomeration



# Sequential execution time

# Sequential execution time

- $\chi$  – time to update element

# Sequential execution time

- $\chi$  – time to update element
- $n$  – number of elements

# Sequential execution time

- $\chi$  – time to update element
- $n$  – number of elements
- $m$  – number of iterations



# Sequential execution time

- $\chi$  – time to update element
- $n$  – number of elements
- $m$  – number of iterations
- Sequential execution time:  $m (n-1) \chi$

# Parallel Execution Time

# Parallel Execution Time

- $\chi$  – time to update element

# Parallel Execution Time

- $\chi$  – time to update element
- $n$  – number of elements

# Parallel Execution Time

- $\chi$  – time to update element
- $n$  – number of elements
- $m$  – number of iterations

# Parallel Execution Time

- $\chi$  – time to update element
- $n$  – number of elements
- $m$  – number of iterations
- $p$  – number of processors

# Parallel Execution Time

- $\chi$  – time to update element
- $n$  – number of elements
- $m$  – number of iterations
- $p$  – number of processors
- $\lambda$  – message latency

# Parallel Execution Time

- $\chi$  – time to update element
- $n$  – number of elements
- $m$  – number of iterations
- $p$  – number of processors
- $\lambda$  – message latency
- Parallel execution time  $m(\chi \lceil (n-1)/p \rceil + 2\lambda)$




# Finding the Maximum Error from measured data

Computed	0.15	0.16	0.16	0.19
Correct	0.15	0.16	0.17	0.18
Error (%)	0.00%	0.00%	6.25%	5.26%


Need to do  
a reduction.

6.25%


# Parallel Reduction Evolution



# Parallel Reduction Evolution




# Parallel Reduction Evolution




# Binomial Trees


# Binomial Trees



# Binomial Trees




# Binomial Trees






# Binomial Trees




# Binomial Trees




Subgraph of hypercube


# Finding Global Sum




# Finding Global Sum




# Finding Global Sum




# Finding Global Sum



# Finding Global Sum




# Finding Global Sum







# Finding Global Sum




# Finding Global Sum



# Finding Global Sum




# Finding Global Sum




Binomial Tree


# Agglomeration




# Agglomeration leads to actual communication



# Agglomeration leads to actual communication






# Agglomeration leads to actual communication






# The n-body Problem




# The n-body Problem




# The n-body Problem




# The n-body Problem




# The n-body Problem






# The n-body Problem



# The n-body Problem



# The n-body Problem





# Partitioning

# Partitioning

- Domain partitioning

# Partitioning

- Domain partitioning
- Assume one task per particle

# Partitioning

- Domain partitioning
- Assume one task per particle
- Task has particle's position, velocity vector

# Partitioning

- Domain partitioning
- Assume one task per particle
- Task has particle's position, velocity vector
- Iteration





# Partitioning

- Domain partitioning
- Assume one task per particle
- Task has particle's position, velocity vector
- Iteration
  - ◆ Get positions of all other particles

# Partitioning


- Domain partitioning
- Assume one task per particle
- Task has particle's position, velocity vector
- Iteration
  - ◆ Get positions of all other particles
  - ◆ Compute new position, velocity

# Gather













# Gather




# All-gather





# All-gather




# Complete Graph for All-gather -- operations shown, no ordering required




# Complete Graph for All-gather -- operations shown, no ordering required



# Hypercube-based All-gather -- ordering required



# Complete Graph for All-gather



# Communication Time

Complete graph


$$(p-1)\left(\lambda + \frac{n/p}{\beta}\right) = (p-1)\lambda + \frac{n(p-1)}{\beta p}$$

Hypercube





$$\sum_{i=1}^{\log p} \left( \lambda + \frac{2^{i-1}n}{\beta p} \right) = \lambda \log p + \frac{n(p-1)}{\beta p}$$




# Adding Data Input



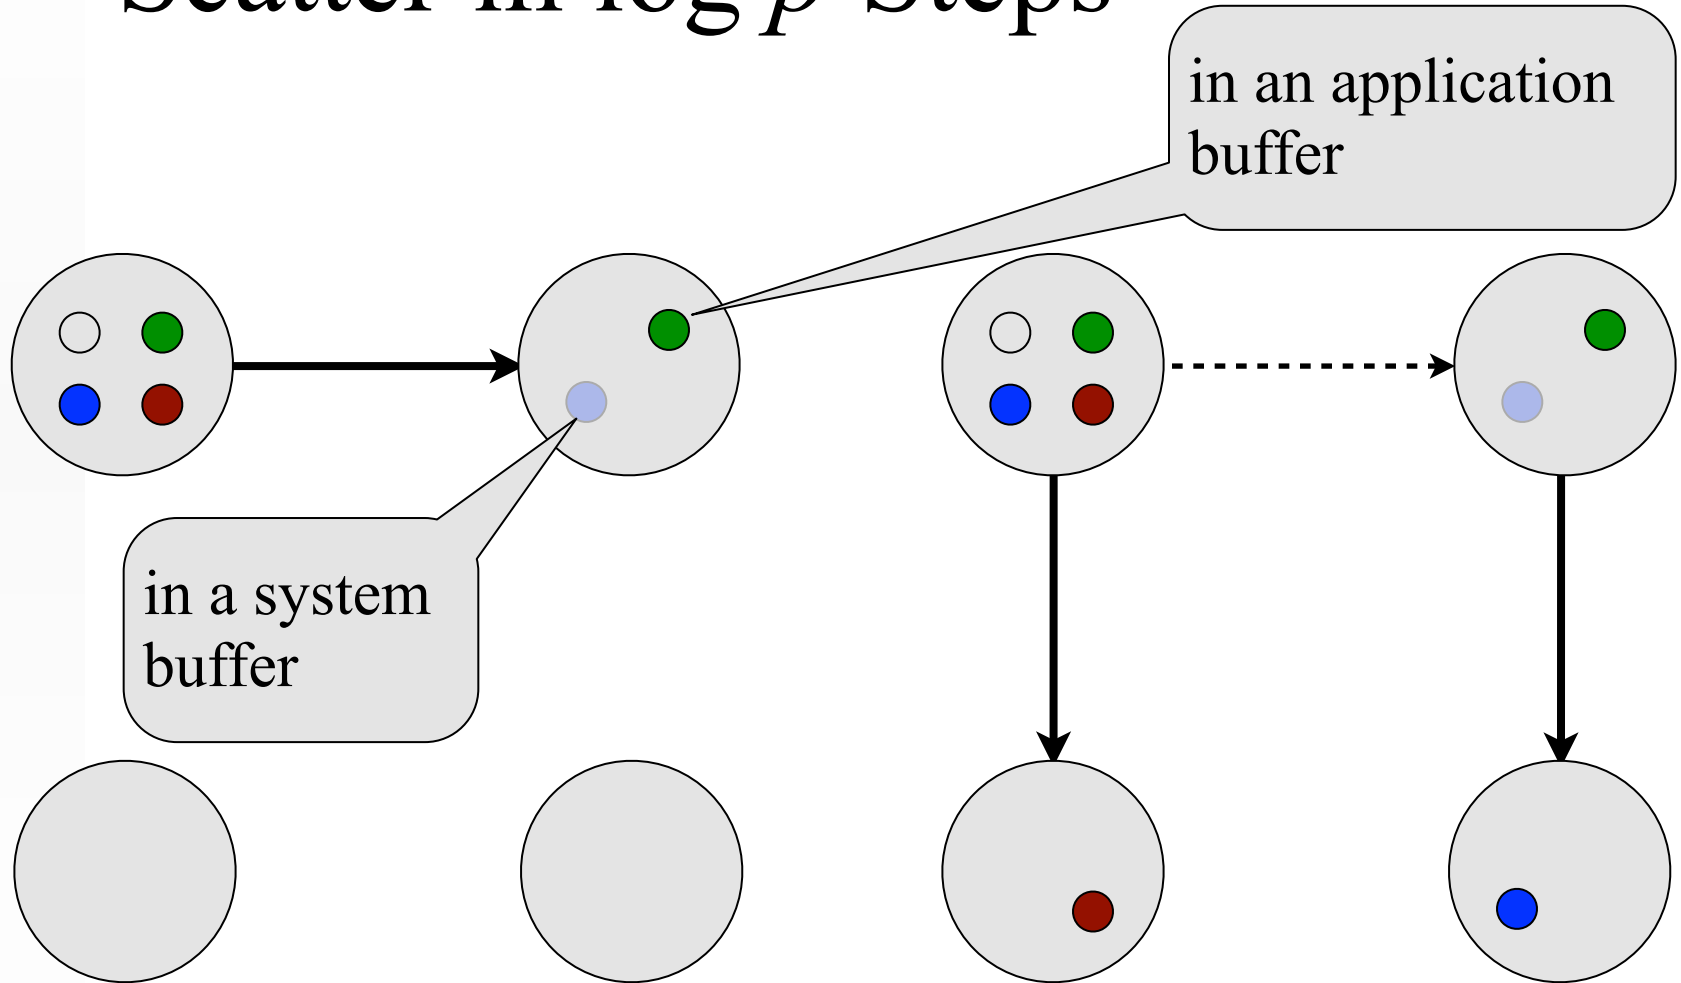
# Scatter



# Scatter



# Scatter in $\log p$ Steps



# Summary: Task/channel Model

# Summary: Task/channel Model

- Parallel computation

# Summary: Task/channel Model

- Parallel computation
  - ◆ Set of tasks

# Summary: Task/channel Model

- Parallel computation
  - ◆ Set of tasks
  - ◆ Interactions through channels



# Summary: Task/channel Model

- Parallel computation
  - ◆ Set of tasks
  - ◆ Interactions through channels
- Good designs

# Summary: Task/channel Model

- Parallel computation
  - ◆ Set of tasks
  - ◆ Interactions through channels
- Good designs
  - ◆ Maximize local computations

# Summary: Task/channel Model

- Parallel computation
  - ◆ Set of tasks
  - ◆ Interactions through channels
- Good designs
  - ◆ Maximize local computations
  - ◆ Minimize communications

# Summary: Task/channel Model

- Parallel computation
  - ◆ Set of tasks
  - ◆ Interactions through channels
- Good designs
  - ◆ Maximize local computations
  - ◆ Minimize communications
  - ◆ Scale up

# Summary: Design Steps

# Summary: Design Steps

- Partition computation

# Summary: Design Steps

- Partition computation
- Agglomerate tasks

# Summary: Design Steps

- Partition computation
- Agglomerate tasks
- Map tasks to processors



# Summary: Design Steps

- Partition computation
- Agglomerate tasks
- Map tasks to processors
- Goals

# Summary: Design Steps

- Partition computation
- Agglomerate tasks
- Map tasks to processors
- Goals
  - ◆ Maximize processor utilization

# Summary: Design Steps

- Partition computation
- Agglomerate tasks
- Map tasks to processors
- Goals
  - ◆ Maximize processor utilization
  - ◆ Minimize inter-processor communication

# Summary: Fundamental Algorithms

# Summary: Fundamental Algorithms

## ■ Reduction

# Summary: Fundamental Algorithms

- Reduction
- Gather and scatter

# Summary: Fundamental Algorithms

- Reduction
- Gather and scatter
- All-gather