
CS250

Computer Architecture

Øæ| 20FG

Part I : More MIPS Instructions and Function Calls

Outline

- We'll go into more detail about the ISA.
 - Pseudo-instructions
 - Using branches for conditionals

Pseudo-instructions

- MIPS assemblers support **pseudo-instructions** that give the illusion of a more expressive instruction set, but are actually translated into one or more simpler, “real” instructions.
- Examples: the **li** and **move** pseudo-instructions:

```
li    $a0, 2000      # Load immediate 2000 into $a0
move  $a1, $t0       # Copy $t0 into $a1
```

- They are probably clearer than their corresponding MIPS instructions:

```
addi  $a0, $0, 2000  # Initialize $a0 to 2000
add   $a1, $t0, $0   # Copy $t0 into $a1
```

- We'll see lots more pseudo-instructions this semester.

— Unless otherwise stated, you can always use pseudo-instructions in your assignments and on exams.

Control flow in high-level languages

- The instructions in a program usually execute one after another, but it's often necessary to alter the normal control flow.
- **Conditional statements** execute only if some test expression is true.

```
// Find the absolute value of *a0
v0 = *a0;
if (v0 < 0)
    v0 = -v0;           // This might not be executed
v1 = v0 + v0;
```

- **Loops** cause some statements to be executed many times.

```
// Sum the elements of a five-element array a0
v0 = 0;
t0 = 0;
while (t0 < 5) {
    v0 = v0 + a0[t0];    // These statements will
    t0++;                // be executed five times
}
```

Control-flow graphs

- It can be useful to draw **control-flow graphs** when writing loops and conditionals in assembly:

```
// Find the absolute value of *a0
v0 = *a0;
if (v0 < 0)
    v0 = -v0;
v1 = v0 + v0;
```

```
// Sum the elements of a0
v0 = 0;
t0 = 0;
while (t0 < 5) {
    v0 = v0 + a0[t0];
    t0++;
}
```

MIPS control instructions

- In section, we introduced some of MIPS's control-flow instructions

`j` // for unconditional jumps
`bne` and `beq` // for conditional branches
`slt` and `slti` // set if less than (w/ and w/o an immediate)

- And how to implement loops
- Today, we'll talk about
 - MIPS's pseudo branches
 - if/else

What does this code do?

```
label:  sub    $a0, $a0, 1  
        bne    $a0, $zero, label
```



Pseudo-branches

- The MIPS processor only supports two branch instructions, **beq** and **bne**, but to simplify your life the assembler provides the following other branches:

```
blt    $t0, $t1, L1 // Branch if $t0 < $t1
ble    $t0, $t1, L2 // Branch if $t0 <= $t1
bgt    $t0, $t1, L3 // Branch if $t0 > $t1
bge    $t0, $t1, L4 // Branch if $t0 >= $t1
```

- There are also immediate versions of these branches, where the second source is a constant instead of a register.
- Later this semester we'll see how supporting just beq and bne simplifies the processor design.

Implementing pseudo-branches

- Most pseudo-branches are implemented using `slt`. For example, a branch-if-less-than instruction `blt $a0, $a1, Label` is translated into the following.

```
slt $at, $a0, $a1    // $at = 1 if $a0 < $a1  
bne $at, $0, Label  // Branch if $at != 0
```

- This supports immediate branches, which are also pseudo-instructions. For example, `blti $a0, 5, Label` is translated into two instructions.


```
slti $at, $a0, 5      // $at = 1 if $a0 < 5  
bne $at, $0, Label  // Branch if $a0 < 5
```

- All of the pseudo-branches need a register to save the result of `slt`, even though it's not needed afterwards.
 - MIPS assemblers use register `$1`, or `$at`, for temporary storage.
 - You should be careful in using `$at` in your own programs, as it may be overwritten by assembler-generated code.

Translating an if-then statement

- We can use branch instructions to translate if-then statements into MIPS assembly code.

```
v0 = *a0;  
if (v0 < 0)  
    v0 = -v0;  
v1 = v0 + v0;
```



- Sometimes it's easier to *invert* the original condition.
 - In this case, we changed “continue if $v0 < 0$ ” to “skip if $v0 \geq 0$ ”.
 - This saves a few instructions in the resulting assembly code.

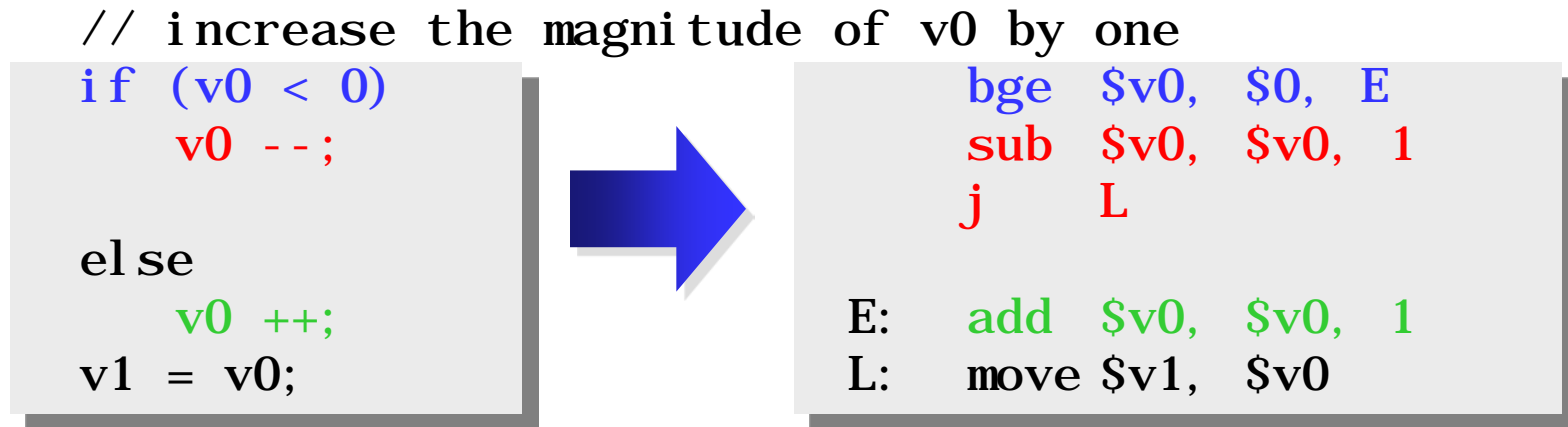


Control-flow Example

- Let's write a program to count how many bits are set in a 32-bit word.

Translating an if-then-else statements


- If there is an **else** clause, it is the target of the conditional branch
 - And the **then** clause needs a jump over the **else** clause



- Dealing with else-if code is similar, but the target of the first branch will be another if statement.
 - Drawing the control-flow graph can help you out.

Functions in MIPS

- We'll talk about the 3 steps in handling function calls:
 1. The program's flow of control must be changed.
 2. Arguments and return values are passed back and forth.
 3. Local variables can be allocated and destroyed.
- And how they are handled in MIPS:
 - New instructions for calling functions.
 - Conventions for sharing registers between functions.
 - Use of a stack.



Control flow in C

- Invoking a function changes the control flow of a program twice.
 1. **Calling** the function
 2. **Returning** from the function
- In this example the **main** function calls **fact** twice, and **fact** returns twice—but to *different* locations in **main**.
- Each time **fact** is called, the CPU has to remember the appropriate **return address**.
- Notice that **main** itself is also a function! It is, in effect, called by the operating system when you run the program.

```
int main()
{
    ...
    t1 = fact(8);
    t2 = fact(3);
    t3 = t1 + t2;
    ...
}

int fact(int n)
{
    int i, f = 1;
    for (i = n; i > 1; i--)
        f = f * i;
    return f;
}
```

Control flow in MIPS

- MIPS uses the jump-and-link instruction **jal** to call functions.
 - The jal saves the return address (the address of the *next* instruction) in the dedicated register **\$ra**, before jumping to the function.
 - jal is the only MIPS instruction that can access the value of the program counter, so it can store the return address PC+4 in \$ra.

jal Fact

- To transfer control back to the caller, the function just has to jump to the address that was stored in \$ra.

jr \$ra

- Let's now add the jal and jr instructions that are necessary for our factorial example.

Data flow in C

- Functions accept **arguments** and produce **return values**.
- The **blue** parts of the program show the actual and formal arguments of the fact function.
- The **purple** parts of the code deal with returning and using a result.

```
int main()
{
    ...
    t1 = fact(8);
    t2 = fact(3);
    t3 = t1 + t2;
    ...
}

int fact(int n)
{
    int i, f = 1;
    for (i = n; i > 1; i--)
        f = f * i;
    return f;
}
```


Data flow in MIPS

- MIPS uses the following conventions for function arguments and results.
 - Up to four function arguments can be “passed” by placing them in argument registers **\$a0-\$a3** before calling the function with jal.
 - A function can “return” up to two values by placing them in registers **\$v0-\$v1**, before returning via jr.
- These conventions are not enforced by the hardware or assembler, but programmers agree to them so functions written by different people can interface with each other.
- Later we’ll talk about handling additional arguments or return values.

A note about types

- Assembly language is **untyped**—there is no distinction between integers, characters, pointers or other kinds of values.
- It is up to ***you*** to “type check” your programs. In particular, make sure your function arguments and return values are used consistently.
- For example, what happens if somebody passes the *address* of an integer (instead of the integer itself) to the fact function?

The big problem so far

- There is a big problem here!
 - The main code uses `$t1` to store the result of `fact(8)`.
 - But `$t1` is also used within the `fact` function!
- The subsequent call to `fact(3)` will overwrite the value of `fact(8)` that was stored in `$t1`.

Nested functions

- A similar situation happens when you call a function that then calls another function.
- Let's say A calls B, which calls C.
 - The arguments for the call to C would be placed in \$a0-\$a3, thus *overwriting* the original arguments for B.
 - Similarly, `jal C` overwrites the return address that was saved in \$ra by the earlier `jal B`.

```
A:    ...  
      # Put B's args in $a0-$a3  
      jal B      # $ra = A2  
A2:   ...
```

```
B:    ...  
      # Put C's args in $a0-$a3,  
      # erasing B's args!  
      jal C      # $ra = B2  
B2:   ...  
      jr $ra     # where does  
                  # this go???
```

```
C:    ...  
      jr $ra
```

Spilling registers

- The CPU has a limited number of registers for use by all functions, and it's possible that several functions will need the same registers.
- We can keep important registers from being overwritten by a function call, by saving them before the function executes, and restoring them after the function completes.
- But there are two important questions.
 - Who is responsible for saving registers—the caller or the callee?
 - Where exactly are the register contents saved?



Who saves the registers?

- Who is responsible for saving important registers across function calls?
 - The caller knows which registers are important to it and should be saved.
 - The callee knows exactly which registers it will use and potentially overwrite.
- However, in the typical “black box” programming approach, the caller and callee do not know anything about each other’s implementation.
 - Different functions may be written by different people or companies.
 - A function should be able to interface with any client, and different implementations of the same function should be substitutable.
- So how can two functions cooperate and share registers when they don’t know anything about each other?

The caller could save the registers...

- One possibility is for the *caller* to save any important registers that it needs before making a function call, and to restore them after.
- But the caller does not know what registers are actually written by the function, so it may save more registers than necessary.
- In the example on the right, *frodo* wants to preserve *\$a0*, *\$a1*, *\$s0* and *\$s1* from *gollum*, but *gollum* may not even use those registers.

```
frodo: li    $a0, 3
        li    $a1, 1
        li    $s0, 4
        li    $s1, 1

        # Save registers
        # $a0, $a1, $s0, $s1

        jal   gollum

        # Restore registers
        # $a0, $a1, $s0, $s1

        add   $v0, $a0, $a1
        add   $v1, $s0, $s1
        jr    $ra
```

...or the callee could save the registers...

- Another possibility is if the *callee* saves and restores any registers it might overwrite.
- For instance, a `gollum` function that uses registers `$a0`, `$a2`, `$s0` and `$s2` could save the original values first, and restore them before returning.
- But the callee does not know what registers are important to the caller, so again it may save more registers than necessary.

```
gollum:
```

```
# Save registers  
# $a0 $a2 $s0 $s2
```

```
li    $a0, 2  
li    $a2, 7  
li    $s0, 1  
li    $s2, 8
```

```
...
```

```
# Restore registers  
# $a0 $a2 $s0 $s2
```

```
jr    $ra
```


...or they could work together

- MIPS uses conventions again to split the register spilling chores.
- The *caller* is responsible for saving and restoring any of the following **caller-saved registers** that it cares about.

\$t0-\$t9

\$a0-\$a3

\$v0-\$v1

In other words, the callee may freely modify these registers, under the assumption that the caller already saved them if necessary.

- The *callee* is responsible for saving and restoring any of the following **callee-saved registers** that it uses. (Remember that \$ra is “used” by jal.)

\$s0-\$s7

\$ra

Thus the caller may assume these registers are not changed by the callee.

— \$ra is tricky; it is saved by a callee who is also a caller.

- Be especially careful when writing nested functions, which act as both a caller and a callee!

Register spilling example

- This convention ensures that the caller and callee together save all of the important registers—frodo only needs to save registers `$a0` and `$a1`, while gollum only has to save registers `$s0` and `$s2`.

```
frodo:  li    $a0, 3
        li    $a1, 1
        li    $s0, 4
        li    $s1, 1

        # Save registers
        # $a0, $a1, $ra

        jal   gollum

        # Restore registers
        # $a0, $a1, $ra

        add   $v0, $a0, $a1
        add   $v1, $s0, $s1
        jr    $ra

gollum:                                     # Save registers
                                           # $s0 and $s2

        li    $a0, 2
        li    $a2, 7
        li    $s0, 1
        li    $s2, 8
        ...

        # Restore registers
        # $s0 and $s2


        jr    $ra
```

How to fix factorial

- In the factorial example, main (the caller) should save two registers.
 - `$t1` must be saved before the second call to fact.
 - `$ra` will be implicitly overwritten by the jal instructions.
- But fact (the callee) does not need to save anything. It only writes to registers `$t0`, `$t1` and `$v0`, which should have been saved by the caller.

Where are the registers saved?

- Now we know who is responsible for saving which registers, but we still need to discuss where those registers are saved.
- It would be nice if each function call had its own private memory area.
 - This would prevent other function calls from overwriting our saved registers—otherwise using memory is no better than using registers.
 - We could use this private memory for other purposes too, like storing local variables.




Function calls and stacks

- Notice function calls and returns occur in a stack-like order: the most recently called function is the first one to return.

1. Someone calls A
2. A calls B
3. B calls C
4. C returns to B
5. B returns to A
6. A returns

- Here, for example, C must return to B *before* B can return to A.




Stacks and function calls

- It's natural to use a **stack** for function call storage. A block of stack space, called a **stack frame**, can be allocated for each function call.
 - When a function is called, it creates a new frame onto the stack, which will be used for local storage.
 - Before the function returns, it must pop its stack frame, to restore the stack to its original state.
- The stack frame can be used for several purposes.
 - Caller- and callee-save registers can be put in the stack.
 - The stack frame can also hold local variables, or extra arguments and return values.



The MIPS stack

- In MIPS machines, part of main memory is reserved for a stack.
 - The stack grows downward in terms of memory addresses.
 - The address of the top element of the stack is stored (by convention) in the “stack pointer” register, $\$sp$.
- MIPS does not provide “push” and “pop” instructions. Instead, they must be done explicitly by the programmer.



Pushing elements


- To **push** elements onto the stack:
 - Move the stack pointer **\$sp** down to make room for the new data.
 - Store the elements into the stack.
- For example, to push registers **\$t1** and **\$t2** onto the stack:

```
sub $sp, $sp, 8  
sw  $t1, 4($sp)  
sw  $t2, 0($sp)
```


- An equivalent sequence is:

```
sw  $t1, -4($sp)  
sw  $t2, -8($sp)  
sub $sp, $sp, 8
```

- Before and after diagrams of the stack are shown on the right.



Before



After

Accessing and popping elements

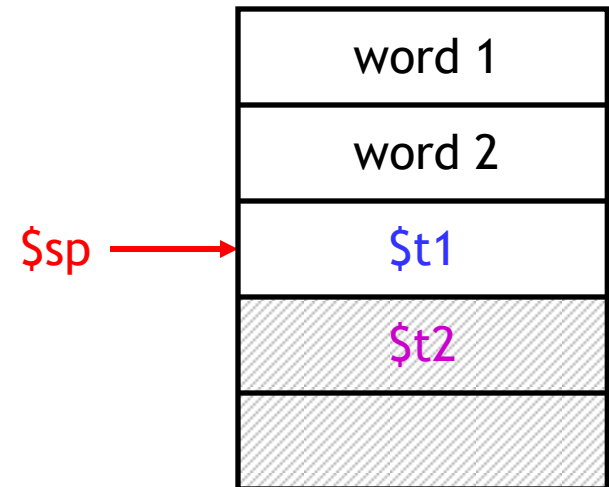

- You can access any element in the stack (not just the top one) if you know where it is relative to `$sp`.
- For example, to retrieve the value of `$t1`:

```
lw    $s0, 4($sp)
```

- You can **pop**, or “erase,” elements simply by adjusting the stack pointer upwards.
- To pop the value of `$t2`, yielding the stack shown at the bottom:

```
addi  $sp, $sp, 4
```


- Note that the popped data is still present in memory, but data past the stack pointer is considered invalid.



Stack

The stack can be used for

- storing return addresses
- storing register values
- parameter passing
- ...



$\$sp = \$sp - 12$

Factorial

- Compute $n!$
- Recall that
 - $0! = 1$
 - $1! = 1$
 - $n! = n(n-1)!$
- Store on the stack
 - $\$s0 = n$, the parameter
 - $\$ra$, the return address

factorial:

```
bgt $a0,$0,gen    # if $a0>0 goto generic case
li $v0, 1          # base case, 0! = 1
jr $ra            # return
```

```
gen:  subi $sp,$sp,8    # create a stack frame
      sw $s0,0($sp)     # store register $s0's value
      sw $ra,4($sp)     # store return address
      move $s0,$a0      # save argument
      subi $a0,$a0,1    # factorial(n-1)
      jal factorial     # v0 = (n-1)!
      mul $v0,$s0,$v0   # n*(n-1)!
      lw $s0,0($sp)     # restore $s0=n
      lw $ra,4($sp)     # restore $ra
      addi $sp,$sp,8    # destruct the stack frame
      jr $ra           # return
```

Fibonacci

$$\text{fib}(0) = 0$$

$$\text{fib}(1) = 1$$

$$\text{fib}(n) = \text{fib}(n-1) + \text{fib}(n-2)$$

0, 1, 1, 2, 3, 5, 8, 13, 21,...

Fibonacci

```
li $a0, 10          # call fib(10)
jal fib             #
move $s0, $v0        # $s0 = fib(10)
```

fib is another recursive procedure with one argument \$a0
need to store argument \$a0, temporary register \$s0 for intermediate
results, and return address \$ra

```

fib:      subi $sp,$sp,12      # save registers on stack
          sw $a0, 0($sp)      # save $a0 = n
          sw $s0, 4($sp)      # save $s0
          sw $ra, 8($sp)      # save return address $ra
          bgt $a0,1, gen      # if n>1 then goto generic case
          move $v0,$a0        # output = input if n=0 or n=1
          j rest              # goto restore registers
gen:      subi $a0,$a0,1      # param = n-1
          jal fib              # compute fib(n-1)
          move $s0,$v0        # save fib(n-1)
          sub $a0,$a0,1      # set param to n-2
          jal fib              # and make recursive call
          add $v0, $v0, $s0    # $v0 = fib(n-2)+fib(n-1)
rest:     lw  $a0, 0($sp)      # restore registers from stack
          lw  $s0, 4($sp)      #
          lw  $ra, 8($sp)      #
          addi $sp, $sp, 12    # decrease the stack size
          jr $ra

```

Summary

- Today we focused on implementing function calls in MIPS.
 - We call functions using `jal`, passing arguments in registers `$a0-$a3`.
 - Functions place results in `$v0-$v1` and return using `jr $ra`.
- Managing resources is an important part of function calls.
 - To keep important data from being overwritten, registers are saved according to conventions for `caller-save` and `callee-save` registers.
 - Each function call uses stack memory for saving registers, storing local variables and passing extra arguments and return values.
- Assembly programmers must follow many conventions. Nothing prevents a rogue program from overwriting registers or stack memory used by some other function.