

Question 1. The answer is “yes”. If the undirected version is connected then it has a spanning tree T . Let (v, u) be any directed edge that appears (in undirected form) on T : Starting with that directed edge (v, u) , keep walking along never-used directed edges of G . Because there are finitely many edges, the process eventually stops at a vertex that has run out of fresh outgoing edges (all have already been used up). But that vertex has to be v , because every other vertex w has had its edges used up in pairs: 1 incoming when we enter it, followed by 1 outgoing when we exit it, and because there are as many incoming as outgoing we cannot “get stuck” at such a w (i.e., we can only get stuck at the v where we started). This cyclic walk that started at v , took edge (v, u) , and eventually ended back at v , has gone through vertices that are clearly all in the same strongly connected component. Therefore v and u are in the same strongly connected component. As the above logic applies to every edge that appears in the spanning tree T , all the vertices of T (hence of G) are in the same strongly connected component.

Question 2.

1. Sort the points by increasing x coordinate, and let L_y be the resulting sequence of y coordinates. Compute the LIS of L_y : The length of this LIS is the α we seek.
2. Sort the points by decreasing x coordinate, and let L_y be the resulting sequence of y coordinates. Compute the LIS of L_y : The length of this LIS is the β we seek.

Alternative solution. Sort the points by increasing x coordinate, and let L_y be the resulting sequence of y coordinates. Replace, in L_y , every y coordinate y_i by $-y_i$. Use the LIS algorithm on the resulting sequence: The length of the LIS is the β we seek.

3. For each point $p_i = (x_i, y_i)$, let S_i be the subset of S whose x coordinates are $\leq x_i$ (hence S_i includes p_i). Let λ_i be the number of points in a largest subset of S_i that contains p_i and is such that all the points in it are comparable (that is, λ_i is the size of a solution in S_i that is constrained to contain p_i as its rightmost point). Note that $\alpha = \max_{1 \leq i \leq n} \lambda_i$. If $\alpha \geq \sqrt{n}$ then $\max\{\alpha, \beta\} \geq \sqrt{n}$ and we are done. So suppose that $\alpha < \sqrt{n}$. For each $t \in \{1, 2, \dots, \alpha\}$ let Θ_t denote the set of p_i s in S whose $\lambda_i = t$. Observe that every pair of distinct points p_i, p_j for which $\lambda_i = \lambda_j$ must be incomparable (because if p_i dominated p_j then we would have $\lambda_i > \lambda_j$). This means that all the points in a set Θ_t are incomparable, and therefore $\beta \geq |\Theta_t|$ for all $1 \leq t \leq \alpha$. Therefore it suffices to show that at least one of the Θ_t 's contains more than \sqrt{n} points. If you view the n points as pigeons, and view each Θ_t as a pigeonhole, then the fact that there are n pigeons and fewer than \sqrt{n} pigeonholes implies that at least one pigeonhole contains more than \sqrt{n} pigeons, completing the proof.

Question 3. To perform an $\text{Increment}(i, j, x)$ we locate the i th leaf (call it l_i) and the j th leaf l_j (in constant time each). Then in $O(h)$ time we find their lowest common ancestor (call it f) and the $O(h)$ internal nodes that are either right siblings of nodes on the f -to- l_i

path or left siblings of nodes on the f -to- l_j path. For each such node v we add x to its $\delta(v)$. A $Decrement(i, j, x)$ is performed similarly except that we subtract x instead of adding it. Correctness of the above follows from the fact that the path between the root and a leaf whose associated item is to be modified, goes through exactly one of the nodes whose δ value was modified.

Question 4. The first idea that comes to mind is to use the convolution method (explained in class) modified in the following way: When setting up the convolution problem for calculating the contribution, to the score vector C , of a particular alphabetical symbol γ , you replace every occurrence of symbol γ and of symbol $\#$ by 1. You do this 4 times, once for every $\gamma \in \{a, b, c, d\}$.

However, adding the above 4 contributions together gives an overestimate of C because you count 4 times the contribution of a $\#$ matching another $\#$. There are two ways to deal with this:

- *Alternative 1:* You do as in the above anyway, and then later you correct for the overestimate by calculating the contribution (call it $C^{(\#)}$) of all $\#$ -to- $\#$ matches (by doing a separate convolution for it in the usual way: You replace every $\#$ by 1 in both T and P and everything else by 0, etc). You then correct by subtracting 3 times $C^{(\#)}$ (because the $\#$ -to- $\#$ matches were counted 4 times). The second alternative (below) is probably better.
- *Alternative 2:* You first use the convolution method without any modification, effectively ignoring the $\#$ symbols by replacing them with zeroes (so you do one convolution for each $\gamma \in \{a, b, c, d\}$). This gives you the effects of matches that do not involve any $\#$ symbols. Then you do 2 separate convolutions to calculate the effects of matches that involve $\#$ symbols: For the first convolution you replace only the $\#$'s in T by 1, and everything in P by 1 (you can actually do this “convolution” in linear time rather than in $O(n \log n)$ time, because all of P was replaced by 1's). For the second convolution you replace everything except the $\#$'s in T by 1, and only the $\#$'s in P by 1. The first convolution captures the effect of matches of $\#$'s in T with anything in P (including $\#$'s), the second one captures the effect of matches of the 4 alphabetical letters in T with $\#$'s in P . Note that there is no double-counting.

Of course, once we have the C vector, the entries of C equal to n correspond to occurrences of P in T .