

**Question 1.** The crucial observation is this: For any leaf  $v$ , a solution that does not include  $v$  can be replaced by another (equally good) solution in which  $v$  has been included and its parent kicked out. This suggests a solution whose high-level description is as follows:

1. Initialize  $S$  to be empty.
2. Repeat the following until the graph becomes empty:
  - (a) Add to  $S$  every vertex  $v$  that is a leaf.
  - (b) Delete from  $T$  the leaves and the parents of leaves. (Note that  $T$  may become disconnected as a result of these deletions, i.e., it can become a forest of trees rather than a single tree.)
3. Output  $S$ .

The following is an algorithm that implements the above idea in  $O(n)$  time by traversing the tree without modifying it, merely marking the nodes that belong to  $S$  during that traversal.

1. Initially none of the vertices is marked as being in  $S$ .
2. A postorder traversal of the tree is done, and at the moment of assigning a postorder number to a node  $v$  a decision is also made on whether to include it in  $S$  or not, according to the following criterion: Unless at least one child of  $v$  has been marked as being in  $S$ ,  $v$  is marked as being in  $S$ . Hence a leaf is in  $S$ , as is a node none of whose children have been marked as being in  $S$  by the postorder traversal.

**Question 2.** For every edge  $(u, v)$ , either both  $u$  and  $v$  are ignored (if one of them is in  $S$ ) or both are included in  $S$ . At least one of them must be in  $\hat{V}$ , because the definition of  $\hat{V}$  requires it. Therefore in the worst case 2 vertices are being added to  $S$  when only one of them is in  $\hat{V}$ . Therefore  $|S| \leq 2|\hat{V}|$ .

**Question 3.** Initially all jobs are marked as *needy*. As the algorithm proceeds, a job that gets a machine assigned to it becomes marked as *not needy*. A machine  $i$  is said to be *compatible* with a job  $J_k$  if  $l_k \leq i \leq r_k$ . The algorithm is then as follows.

1. Go through the machines in the order  $1, 2, \dots, m$  and, for each such  $i$  do the following:
  - (a) Compute the set of jobs (call it  $S_i$ ) that are still needy, and are compatible with machine  $i$ .
  - (b) From the set  $S_i$ , pick the job  $J_k$  that has the smallest  $r_k$ , assign machine  $i$  to  $J_k$ , and mark  $J_k$  as being not needy. (Of course if  $S_i$  is empty then there is no such  $J_k$  and machine  $i$  remains unused.)

The intuitive rationale for the above greedy choice of  $J_k$  is that, among the jobs in  $S_i$ ,  $J_k$  is the job that is most at risk of remaining permanently needy (once the machine number exceeds its  $r_k$ ).

**Question 4.** For every leaf  $v$ ,  $With[v] = w[v]$  and  $Without[v] = 0$ . For every non-leaf  $v$ , the following holds:

$$With[v] = w[v] + \sum_{x \in L[v]} Without[x]$$

$$Without[v] = \sum_{x \in L[v]} \max\{With[x], Without[x]\}$$

The above suggests computing the *With* and *Without* vectors in a postorder traversal of the tree, because once we have the *With* and *Without* values for the children of a node  $v$ , we can compute them for the node  $v$  in  $O(|L[v]|)$  time, for a total time of  $\sum_v |L[v]| = O(n)$ .