

Question 1. (20 points) Let G be an n -vertex directed graph that has the following properties:

- For every vertex v , the number of other vertices that appear on the adjacency list of v is exactly the same as the number of times that v appears in the adjacency lists of other vertices (in other words, for every vertex v the number of directed edges that have v as head is equal to the number of directed edges that have v as tail).
- The undirected version of G (the undirected graph obtained from G by ignoring edge directions) is connected.

Does G have to be strongly connected, or can it be otherwise? Justify your answer by providing a proof if your answer is “Yes”, giving a counterexample if your answer is “No”.

Question 2. (30 points) Suppose you are given a set S of n points $(x_1, y_1), \dots, (x_n, y_n)$ where the x_i s and y_i s are distinct (i.e., no two are equal). We use p_i as a shorthand for the point (x_i, y_i) . A point p_i is said to *dominate* another point p_j if $x_j < x_i$ and $y_j < y_i$. Two points are *comparable* if one of them dominates the other, and are *incomparable* if neither of them dominates the other. For example, the point $(9.2, 3.3)$ dominates the point $(7.1, 1.2)$, but the two points $(9.2, 3.3)$ and $(4.5, 6.8)$ are incomparable. Let α be the number of points of a largest subset of S in which all the points are pairwise comparable. Let β be the number of points of a largest subset of S in which all the points are pairwise incomparable.

1. Give an $O(n \log n)$ time algorithm for computing α by making use of the longest increasing subsequence (LIS) algorithm we covered in class.
2. Repeat the above for computing β .
3. Prove that $\max\{\alpha, \beta\} \geq \sqrt{n}$. (*Hint:* Use the pigeonhole principle.)

Question 3. (20 points) Let T be a (not necessarily complete or balanced) n -leaf binary tree, of height h , whose leaves initially contain n data items d_1, \dots, d_n (not in sorted order); the i th leftmost leaf initially contains d_i . Assume that h is much smaller than n . We would like to support the following operations, in $O(h)$ time per operation.

1. *Increment*(i, j, x) where $1 \leq i < j \leq n$: Adds (in the sense of arithmetic addition) x to the value of the item d_k associated with the k th leftmost leaf, for all k such that $i < k < j$.
2. *Decrement*(i, j, x) where $1 \leq i < j \leq n$: Subtracts x from the value of the item d_k associated with the k th leftmost leaf, for all k such that $i < k < j$.
3. *Value*(i): Returns the current value d_i associated with the i th leftmost leaf. Even though such an operation can be done in constant time in the initial tree T (i.e., before there have been any *Increment* or *Decrement* operation), by indexing into the i th leaf and reading the d_i value in it, this will no longer be a constant-time operation after there have been many *Increment* and *Decrement* operations.

Note that $Increment(i, j, x)$ and $Decrement(i, j, x)$ do not return anything: Their only effect is on later $Value(i)$ operations (which are the only operations that return a value to the outside world). This is why, even though the number of values affected by these two operations could be proportional to n , it is possible to process them in $O(h)$ time. This is done by storing at each node v a field $\delta(v)$ that is initially zero if v is not a leaf. If v is a leaf then $\delta(v)$ is initialized to be the item stored at that leaf. We assume an array is available whose i th entry points to the i th leaf (so that the i th leaf can be accessed in constant time).

- Explain in detail how $Increment(i, j, x)$ and $Decrement(i, j, x)$ are implemented, in $O(h)$ time, so as to maintain the following invariant: “ $Value(i)$ equals the sum of all the $\delta(v)$ values on the path between the root and the i th leftmost leaf”.

Note that the above implies that a query $Value(i)$ is performed in $O(h)$, time by adding all the $\delta(v)$ values on the path between the root and the i th leaf.

Question 4. (30 points) In this problem we consider the exact pattern matching problem when the alphabet consists of the 5 symbols $\{a, b, c, d, \#\}$ where the special symbol $\#$ matches any symbol (including itself). For example, if $T = ab\#aca\#ab\#a$ and $P = da\#ac$ then P occurs starting at position 3 in T . Give an $O(n \log n)$ time algorithm for determining whether a pattern P of length n occurs in a text T of length $2n$, assuming that $\#$ symbols can occur (possibly $O(n)$ times) in both T and P .

Hint. Use convolution, and beware of double-counting.

Date due: November 21, 2013