

TCP congestion control

Recall:

$$\text{EffectiveWindow} = \text{MaxWindow} - (\text{LastByteSent} - \text{LastByteAcked})$$

where

$$\text{MaxWindow} = \min\{\text{AdvertisedWindow}, \text{CongestionWindow}\}$$

Key question: how to set **CongestionWindow** which, in turn, affects ARQ's sending rate?

- linear increase/exponential decrease
- AIMD
- method B

TCP congestion control components:

(i) Congestion avoidance

→ linear increase/exponential decrease

→ additive increase/exponential decrease (AIMD)

As in Method B, increase `CongestionWindow` linearly,
but decrease exponentially

Upon receiving ACK:


$$\text{CongestionWindow} \leftarrow \text{CongestionWindow} + 1$$

Upon timeout:

$$\text{CongestionWindow} \leftarrow \text{CongestionWindow} / 2$$


But is it correct...

“Linear increase” time diagram:



→ results in exponential increase

What we want:



→ increase by 1 every window

Thus, linear increase update:

$$\begin{aligned} \text{CongestionWindow} &\leftarrow \text{CongestionWindow} \\ &\quad + (1 / \text{CongestionWindow}) \end{aligned}$$

Upon timeout and exponential backoff,

$$\text{SlowStartThreshold} \leftarrow \text{CongestionWindow} / 2$$

(ii) Slow Start

Reset `CongestionWindow` to 1

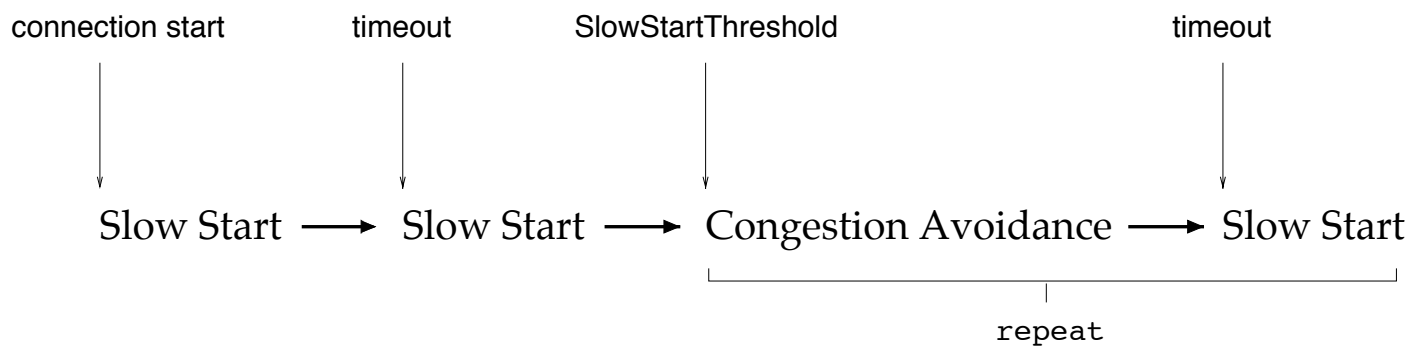
Perform exponential increase

$\text{CongestionWindow} \leftarrow \text{CongestionWindow} + 1$

- Until timeout at start of connection
 - rapidly probe for available bandwidth
- Until `CongestionWindow` hits `SlowStartThreshold` following Congestion Avoidance
 - rapidly climb to safe level
 - “slow” is a misnomer
 - exponential increase is super-fast

Basic dynamics:


- after connection set-up
- before connection tear-down



- most TCP transfers are small
- small files “dominate” Internet TCP connections
- most TCP flows don’t escape **Slow Start**

CongestionWindow evolution:

→ relevant for larger flows



(iii) Exponential timer backoff

$\text{TimeOut} \leftarrow 2 \cdot \text{TimeOut}$ if retransmit

(iv) Fast Retransmit

Upon receiving three duplicate ACKs:

- Transmit next expected segment
 - segment indicated by ACK value
- Perform exponential backoff and commence Slow Start
 - three duplicate ACKs: likely segment is lost
 - react before timeout occurs

TCP Tahoe: features (i)-(iv)

(v) Fast Recovery

Upon Fast Retransmit:

- Skip Slow Start and commence Congestion Avoidance
→ dup ACKs: likely spurious loss
- Insert “inflationary” phase just before Congestion Avoidance

Given sawtooth behavior of TCP's linear increase/exponential backoff:


Why use exponential backoff and not Method D?

- For multimedia streaming (e.g., pseudo real-time), AIMD (Method B) is not appropriate
→ use Method D
- For unimodal case—throughput decreases when system load is excessive—story is more complicated
→ asymmetry in control law needed for stability

Congestion control and selfishness

- to be or not to be selfish ...
- John von Neumann, John Nash, ...

Ex.: “tragedy of commons,” Garrett Hardin, '68



- if everyone acts selfishly, no one wins
 - in fact, everyone loses
- can this be prevented?

Ex.: Prisoner's Dilemma game

→ formalized by Tucker in 1950

→ “cold war”

- both cooperate (i.e., stay mum): 1 year each
- both selfish (i.e., rat on the other): 5 years each
- one cooperative/one selfish: 9 vs. 0 years

		<i>Bob</i>	
		C	N
<i>Alice</i>	C	5, 5	1, 9
	N	9, 1	3, 3

→ payoff matrix

→ what would “rational” prisoners do?

When cast as congestion control game:

		<i>Bob</i>	
		C	N
<i>Alice</i>	C	5, 5	1, 9
	N	9, 1	3, 3



- Alice and Bob share network bandwidth
- (a, b) : throughput (Mbps) achieved by Alice/Bob
- upon congestion: back off or escalate?
- equivalent to Prisoner's dilemma

Rational: in the sense of seeking selfish gain

- both choose strategy “N”
- called Nash equilibrium
- note: stable state
- why: strategy “N” dominates strategy “C”

5 regular (cooperative) TCP flows:

→ share 11 Mbps WLAN bottleneck link



4 regular (cooperative) TCP flows and 1 noncooperative TCP flow:

→ same benchmark set-up

