

Route or path: criteria of goodness

- Hop count
- Delay
- Bandwidth
- Loss rate

Composition of goodness metric:

—→ quality of end-to-end path

- Additive: hop count, delay
- Min: bandwidth
- Multiplicative: loss rate

Goodness of routing:

- assume N users or sessions
- suppose path metric is delay

Two approaches:

- system optimal routing
 - choose paths to minimize $\frac{1}{N} \sum_{i=1}^N D_i$
 - good for the system as a whole
- user optimal routing
 - each user i chooses path to minimize D_i
 - selfish route selections by each user
 - end result may not be good for system as a whole

Algorithms

Find short, in particular, shortest paths from source to destination.

Key observation on shortest paths:

- Assume p is a shortest path from S to D

$$\rightarrow S \overset{p}{\rightsquigarrow} D$$



- Pick any intermediate node X on the path
- Consider the two segments p_1 and p_2

$$\rightarrow S \overset{p_1}{\rightsquigarrow} X \overset{p_2}{\rightsquigarrow} D$$

- The path p_1 from S to X is a shortest path, and so is the path p_2 from X to D

→ leads to Dijkstra's algorithm

Illustration:



→ suggests algorithm for finding shortest path

Leads to Dijkstra's shortest-path algorithm:

→ single-source all-destination

Features:

- running time: $O(n^2)$ time complexity
 - n : number of nodes
- if heap is used: $O(|E| \log |V|)$
 - $O(n \log n)$ if $|E| = O(n)$
- can also be run “backwards”
 - start from destination D and go to all sources
 - a variant used in inter-domain routing
 - forward version: used in intra-domain routing
- source S requires global link distance knowledge
 - centralized algorithm (center: source S)
 - every router runs Dijkstra with itself as source
 - lots of broadcast management packets

- Internet protocol implementation
 - OSPF (Open Shortest Path First)
 - also called link state algorithm
 - broadcast protocol
- builds minimum spanning tree rooted at S :
 - to all destinations
 - if select destination: called multicasting
 - multicast group
 - standardized feature of IETF but not actively utilized on Internet
 - complexity including group membership management

Distributed/decentralized shortest path algorithm:

- Bellman-Ford algorithm
- based on shortest path decomposition property

Key procedure:

- Each node X maintains current shortest distance to all other nodes
 - a distance vector
- Each node advertises to neighbors its current best distance estimates
 - i.e., neighbors exchange distance vectors
- Each node updates shortest paths based on neighbors' advertised information
 - same update criterion as Dijkstra's algorithm

Features:

- running time: $O(n^3)$
- each source or router only talks to neighbors
 - local interaction
 - no need to send update if no change
 - if change, entire distance vector must be sent
- knows shortest distance, but not path
 - just the next hop is known
- elegant but additional issues compared to Dijkstra's algorithm
 - e.g., stability
- Internet protocol implementation
 - RIP (Routing Information Protocol)

QoS routing:

Given two or more performance metrics—e.g., delay and bandwidth—find path with delay less than target delay D (e.g., 100 ms) and bandwidth greater than target bandwidth B (e.g., 1.5 Mbps)

- from shortest path to best QoS path
- multi-dimensional QoS metric
- other: jitter, hop count, etc.

How to find best QoS path that satisfies all requirements?

Brute-force

- enumerate all possible paths
- rank them

Policy routing:


- meaning of “policy” is not precisely defined
- almost anything goes

Criteria include:

- Performance
 - e.g., short paths
- Trust
 - what is “trust”?
- Economics
 - pricing
- Politics, etc.

BGP (Border Gateway Protocol):

→ inter-domain routing



→ “peering” between two domains

→ typical: customer-provider relationship

→ in some cases: equals (true peers)

→ Internet exchanges: multiple domains meet up

- CIDR addressing

- i.e., $a.b.c.d/x$

- Purdue: 128.10.0.0/16, 128.210.0.0/16, 204.52.32.0/20

- check at www.iana.org (e.g., ARIN for US)

- Metric: policy

- e.g., shortest-path, trust, pricing

- meaning of “shortest”: delay, router hop, AS hop

- mechanism: path vector routing

- BGP update message

BGP route update:

—→ BGP update message propagation

BGP update message format:

$$\text{ASNA}_k \rightarrow \cdots \rightarrow \text{ASNA}_2 \rightarrow \text{ASNA}_1; \text{a.b.c.d/x}$$

Meaning: ASN A_1 (with CIDR address a.b.c.d/x) can be reached through indicated path

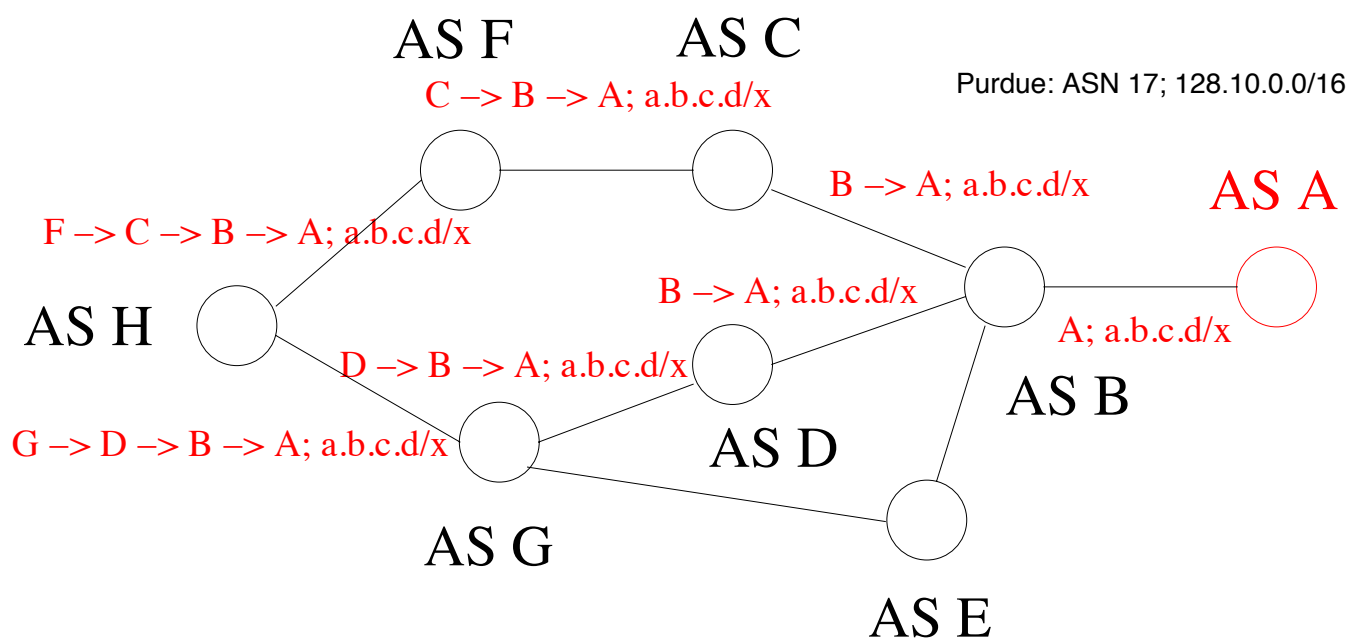

—→ called path vector

—→ also AS-PATH

Some AS numbers:

- Purdue: 17
- BBN: 1
- UUNET: 701
- Level3: 3356
- Abilene (aka “Internet2”): 11537
- AT&T: 7018

Example:



Performance

Route update frequency:

- routing table stability vs. responsiveness
- rule: not too frequently
- 30 seconds
- stability wins
- hard lesson learned from the past (sub-second)
- legacy: TTL

Other factors for route instability:

- selfishness (e.g., fluttering)
- BGP's vector path routing: inherently unstable
- more common: slow convergence
- target of denial-of-service (DoS) attack

Route amplification:

- shortest AS path \neq shortest router path
- e.g., may be several router hops longer
- AS graph vs. router graph
- policy: company in Denmark

Route asymmetry:

- routes are not symmetric
- estimate: $> 50\%$
- mainly artifact of inter-domain policy routing
- various performance implications
- source traceback

Black holes:

- persistent unreachable destination prefixes
- BGP routing problems
- further aggravated by DNS