

Chapter 5

Interlude: μ Scheme in ML

Contents

5.1	Environments	214
5.2	Abstract syntax and values	215
5.3	Evaluation	217
5.4	Primitives	219
5.5	Evaluating definitions	221
5.6	The read-eval-print loop	222
5.7	Initializing and running the interpreter	223
5.8	Building and exporting a program	223
5.9	Free and bound variables	225
5.10	Exercises	227

The interpreters in Chapters 2 through 4 are written in C. C is a very good language in which to write a garbage collector: it is simple and widely used, and it provides the low-level, unsafe features that we need. For illustrating programming languages, however, C is far from ideal. In this and succeeding chapters, we therefore present interpreters written in the functional language Standard ML. Standard ML is particularly well suited to symbolic computing, and the connection between language design, formal semantics, and implementations is much more clearly illustrated by an ML program than by a C program. As an introduction to writing interpreters in ML, and to show the continuity between the first and second halves of this book, we present a μ Scheme interpreter written in ML.

If you happen to be a seasoned ML programmer, you may find our code a bit strange, because we assume you have never seen Standard ML modules. Our code is not broken into modules; it is packed into a single source file. Avoiding modules is poor style, and it makes it impossible to take full advantage of the libraries that come with many implementations of ML, including Standard ML of New Jersey and Moscow ML. But by avoiding modules, we ease the transition for readers whose only previous experience with functional languages is their work with μ Scheme in Chapter 3.

Because we don't use ML modules, we have no formal way to talk about interfaces and to distinguish interfaces from implementations. We work around this problem using a literate-programming trick: we put the types of functions and values, which is mostly what ML interfaces describe, in boxes preceding the implementations. This technique makes it possible to present an interface formally just before we present its implementation. The Noweb processor ensures that the material in the boxes is checked by the ML compiler.

The interpreter in this chapter has the same structure as the interpreter in Chapter 3; as before, we have environments, abstract syntax, an evaluator for expressions, primitives, and an evaluator for definitions. The details of initialization and creating a program are somewhat different in ML than in C, and we use the spelling and capitalization conventions recommended by the SML'97 Standard Basis Library (Gansner and Reppy 2002).

5.1 Environments

In our C implementation, an environment binds names to locations. In our ML implementation, an environment binds names to mutable cells, but that's not all. We actually provide a *polymorphic* implementation in which an environment of type 'a env binds names to values of type 'a. The *type variable* 'a stands for an unknown type. A type variable can be *instantiated* at any type; to get an environment binding names to mutable cells, we instantiate our implementation using 'a = value ref.

We represent an environment of type 'a env as a list of (name, 'a) pairs. The declarations in the box give the interface to our implementation; through some Noweb hackery, they are checked by the ML compiler.

```

214  (environments 214)≡ (224)
    type name = string
    type 'a env = (name * 'a) list
    val emptyEnv = []

    (* lookup and assignment of existing bindings *)
    exception NotFound of name
    fun find (name, []) = raise NotFound name
      | find (name, (n, v)::tail) = if name = n then v else find (name, tail)

    (* adding new bindings *)
    exception BindListLength
    fun bind (name, v, rho) = (name, v) :: rho
    fun bindList (n::vars, v::vals, rho) = bindList (vars, vals, bind (n, v, rho))
      | bindList ([], [], rho) = rho
      | bindList _ = raise BindListLength

```

Because ML strings are immutable, we can use them to represent names directly. We also use exceptions, not an error procedure, to indicate when things have gone wrong. The exceptions we use are listed in Table 5.1.

Exceptions raised at parse time

SyntaxError	Something else went wrong during parsing, i.e., during the execution of readdef.
-------------	--

Exceptions raised at run time

NotFound	A name was looked up in an environment but not found there.
BindListLength	A call to bindList tried to extend an environment, but it passed two lists (names and values) of different lengths.
RuntimeError	Something else went wrong during evaluation, i.e., during the execution of eval.

Table 5.1: Exceptions defined especially for this interpreter

5.2 Abstract syntax and values

An abstract-syntax tree can contain a literal value. A value, if it is a closure, can contain an abstract-syntax tree. These two types are therefore mutually recursive, and in ML, two mutually recursive types have to be declared together, using `and`.¹

```

215 <abstract syntax and values 215>≡ (224) 216a>
    datatype exp = LITERAL of value
                | VAR      of name
                | SET      of name * exp
                | IFX      of exp * exp * exp
                | WHILEX   of exp * exp
                | BEGIN    of exp list
                | APPLY    of exp * exp list
                | LETX     of let_kind * (name * exp) list * exp
                | LAMBDA   of lambda
    and let_kind = LET | LETREC | LETSTAR
    and value = NIL
                | BOOL     of bool
                | NUM      of int
                | SYM      of name
                | PAIR     of value * value
                | CLOSURE   of lambda * value ref env
                | PRIMITIVE of primitive
    withtype primitive = value list -> value (* raises RuntimeError *)
           and lambda  = name list * exp

    exception RuntimeError of string (* error message *)

```

type env 214
 type name 214

These definitions show one refinement of our C implementation; we represent a primitive as an ML function of type `value list -> value`. This simple representation requires no names or tags, and it is suggested by a principle of functional programming: *don't encode a function as data if you can use the function itself*. When we apply a primitive function and something goes wrong, we raise the `RuntimeError` exception.

¹ML is more restrictive than C, in which we can use incomplete structures to distribute the definitions of mutually recursive types over distant source locations, even in different files. Similar results can be achieved in ML by using "two-level types" (Sheard 2001), but the details are beyond the scope of this book.

Definitions are straightforward.

216a $\langle \text{abstract syntax and values 215} \rangle + \equiv$ (224) $\langle 215$

```

datatype def = VAL      of name * exp
              | EXP      of exp
              | DEFINE of name * lambda
              | USE       of name

```

We provide some convenience functions on values. In particular, we show how to embed an ML Boolean or list into its μ Scheme version and how to render a value as a string.

Embedding and projection An S-expression can represent an integer, Boolean, name, function, list, etc. We may sometimes have an ML Boolean, list, or function that we wish to represent as an S-expression, or similarly, an S-expression that we wish to represent as a value of type `bool`. Here we define mappings between type `value` and some other ML types. Because the set of values representable by an ML value of type `value` strictly contains each of the sets of values representable by these ML types, these mappings are called *embedding* and *projection*. Because the `value` type is strictly larger than these ML types, no embedding operation ever fails, but a projection operation might.² For example, although *any* ML function of type `value \rightarrow bool` can be embedded into `value` by using the `PRIMITIVE` constructor, there are values of type `value` that cannot be projected into an ML function of type `value \rightarrow bool`.

Lists and Booleans are straightforward. Function `embedPredicate` is not a true embedding; it takes any function returning `bool` and returns a corresponding function returning `value`. It really embeds the function's result, not the function itself.

$\langle \text{values 216b} \rangle \equiv$ (224) 217a

<code>embedList</code>	: <code>value list \rightarrow value</code>
<code>embedPredicate</code>	: <code>('a \rightarrow bool) \rightarrow ('a \rightarrow value)</code>
<code>bool</code>	: <code>value \rightarrow bool</code>

```

fun embedList []      = NIL
  | embedList (h::t) = PAIR (h, embedList t)
fun embedPredicate f = fn x => BOOL (f x)
fun bool (BOOL b) = b
  | bool _       = true

```

Function `bool` is the projection function, mapping μ Scheme values into ML Booleans.

²This property is a general characteristic of any embedding/projection pair. Mathematical terminology may clarify; an embedding e of S into S' is an injection from $S \rightarrow S'$. The corresponding projection π_e is a left inverse of the embedding; that is $\pi_e \circ e$ is the identity function on S . There is no corresponding guarantee for $e \circ \pi_e$; for example, π_e may be undefined (\perp) on some elements of S' , or $e(\pi_e(x))$ may not equal x .

Printing We render an S-expression as a string.

217a *<values 216b>+≡* (224) <216b

```

fun valueString (NIL)      = "()"
  | valueString (BOOL b)   = if b then "#t" else "#f"
  | valueString (NUM n)    = String.map (fn #"~" => #"- " | c => c) (Int.toString n)
  | valueString (SYM v)    = v
  | valueString (PAIR (car, cdr)) =
    let fun tail (PAIR (car, cdr)) = " " ^ valueString car ^ tail cdr
        | tail NIL = ""
        | tail v = " . " ^ valueString v ^ ""
    in "(" ^ valueString car ^ tail cdr
    end
  | valueString (CLOSURE _) = "<procedure>"
  | valueString (PRIMITIVE _) = "<procedure>"

```

valueString : value -> string

The syntax `Int.toString` indicates the `toString` function from the standard module `Int`. This function, which is part of ML's Standard Basis Library, converts an integer to a string. We use another standard function, `String.map`, to change the minus sign from the ML convention (`~`) to the Scheme convention (`-`).

5.3 Evaluation

The machinery above is enough to write the evaluator, which takes an expression and an environment and produces a value. To make the evaluator easy to write, we do most of the work of evaluation in the nested function `ev`, which inherits the environment `rho` from the outer function `eval`. Because most AST nodes are evaluated in the same environment as their parents, we can evaluate most of them by calling `ev`, which lets `rho` be implicit.

217b *<evaluation 217b>≡* (224) 219b>

```

<definition of separate 218d>
fun eval (e, rho) =
  let fun ev (LITERAL n) = n
      <more alternatives for ev 217c>
  in ev e
  end

```

eval : exp * value ref env -> value

Because `rho` binds names to mutable ref cells, not to values directly, we need the ML functions `!` and `:=` to read and change the contents of the ref cells. The right-hand side of `SET` is evaluated in the same environment as the `SET`, so we use `ev`.

217c *<more alternatives for ev 217c>≡* (217b) 218a>

```

| ev (VAR v) = !(find(v, rho))
| ev (SET (n, e)) =
  let val v = ev e
  in find (n, rho) := v;
  v
  end

```

BOOL,	
in Typed	
μScheme	268a
in μML	287c
in μScheme	215
CLOSURE,	
in Typed	
μScheme	268a
in μML	287c
in μScheme	215
find	214
LITERAL	215
NIL,	
in Typed	
μScheme	268a
in μML	287c
in μScheme	215
NUM,	
in Typed	
μScheme	268a
in μML	287c
in μScheme	215
PAIR,	
in Typed	
μScheme	268a
in μML	287c
in μScheme	215
PRIMITIVE,	
in Typed	
μScheme	268a
in μML	287c
in μScheme	215
SET	215
SYM,	
in Typed	
μScheme	268a
in μML	287c
in μScheme	215
VAR	215

Using the projection function `bool` makes it easy to map μ Scheme control flow onto ML control-flow constructs.

```

218a  <more alternatives for ev 217c>+≡ (217b) <217c 218b>
      | ev (IFX (e1, e2, e3)) = ev (if bool (ev e1) then e2 else e3)
      | ev (WHILEX (guard, body)) =
          if bool (ev guard) then
              (ev body; ev (WHILEX (guard, body)))
          else
              NIL
      | ev (BEGIN es) =
          let fun b (e::es, lastval) = b (es, ev e)
              | b ( [], lastval) = lastval
          in b (es, BOOL false)
          end

```

Capturing a closure is as simple as in C. Applying a primitive function is even simpler; we just evaluate the arguments and apply the primitive to the results.

```

218b  <more alternatives for ev 217c>+≡ (217b) <218a 219a>
      | ev (LAMBDA l) = CLOSURE (l, rho)
      | ev (APPLY (f, args)) =
          (case ev f
           of PRIMITIVE prim => prim (map ev args)
            | CLOSURE clo    => <apply closure clo to args 218c>
            | v => raise RuntimeError ("Applied non-function " ^ valueString v)
           )

```

Applying a closure is more interesting. To apply a μ Scheme closure correctly, we have to create fresh locations to hold the values of the actual parameters. In C, we wrote the function `allocate` for this purpose; in ML, the built-in function `ref` does the same thing: create a new location and initialize its contents with a value.

```

APPLY      215
args       687c
BEGIN      215
bindList   214
BindListLength
214
BOOL       215
bool       216b
clo        687c
CLOSURE    215
ev,
  in Typed
     $\mu$ Scheme
      686c
  in  $\mu$ Scheme
    217b
eval,
  in Typed
     $\mu$ Scheme
      686c
  in  $\mu$ Scheme
    217b
IFX        215
LAMBDA     215
NIL        215
PRIMITIVE  215
rho        217b
RuntimeError,
  in Typed
     $\mu$ Scheme
      268a
  in  $\mu$ Scheme
    215
valueString 217a
WHILEX     215

```

```

<apply closure clo to args 218c>≡ (218b)
  let val ((formals, body), savedrho) = clo
      val actuals = map ev args
  in eval (body, bindList (formals, map ref actuals, savedrho))
      handle BindListLength =>
          raise RuntimeError ("Wrong number of arguments to closure; " ^
                              "expected (" ^ spaceSep formals ^ ")")
  end

```

Function `separate` helps print a readable list of formals for the error message. Function `spaceSep` is a common special case.

```

<definition of separate 218d>≡ (217b)
  fun separate (zero, sep) = (* print list with separator *)
      let fun s []      = zero
          | s [x]       = x
          | s (h::t)    = h ^ sep ^ s t
      in s
      end
  val spaceSep = separate ("", " ") (* print separated by spaces *)

```

To interpret `let`, it is easiest to unzip the list of pairs `bs` into a pair of lists (`names`, `values`). For `let*`, however, it is easier to walk the bindings one pair at a time. The implementation of `letrec` does both.

The function `ListPair.unzip` is from the `ListPair` module in the Standard Basis Library.

```
219a <more alternatives for ev 217c>+≡ (217b) <218b
ListPair.unzip : ('a * 'b) list -> 'a list * 'b list
| ev (LETX (LET, bs, body)) =
  let val (names, values) = ListPair.unzip bs
  in eval (body, bindList (names, map (ref o ev) values, rho))
  end
| ev (LETX (LETSTAR, bs, body)) =
  let fun step ((n, e), rho) = bind (n, ref (eval (e, rho)), rho)
  in eval (body, foldl step rho bs)
  end
| ev (LETX (LETREC, bs, body)) =
  let val (names, values) = ListPair.unzip bs
      val rho' = bindList (names, map (fn _ => ref NIL) values, rho)
      val bs = map (fn (n, e) => (n, eval (e, rho'))) bs
  in List.app (fn (n, v) => find (n, rho') := v) bs;
      eval (body, rho')
  end
```

5.4 Primitives

Here are the primitives. All are either binary or unary operators. As in C, we would like to reuse the code that does the arity checks. In ML, it's easy; we use higher-order functions to write the arity checks just once. If a check fails, function `arityError` raises `RuntimeError` with a suitable message.

```
219b <evaluation 217b>+≡ (224) <217b 219c>
unaryOp  : (value          -> value) -> (value list -> value)
binaryOp : (value * value -> value) -> (value list -> value)

fun arityError n args =
  raise RuntimeError ("primitive function expected " ^ Int.toString n ^
    "arguments; got " ^ Int.toString (length args))
fun binaryOp f = (fn [a, b] => f (a, b) | args => arityError 2 args)
fun unaryOp f = (fn [a]    => f a    | args => arityError 1 args)
```

<code>bind</code>	214
<code>bindList</code>	214
<code>ev</code>	217b
<code>eval</code>	217b
<code>find</code>	214
<code>LET</code>	215
<code>LETREC</code>	215
<code>LETSTAR</code>	215
<code>LETX</code>	215
<code>NIL</code>	215
<code>NUM</code>	215
<code>rho</code>	217b
<code>RuntimeError</code>	215

By using higher-order functions, we encapsulate the ideas of “binary operator” and “unary operator” in a general way. As we subdivide our primitives into arithmetic, predicates, list primitives, and other, we use more higher-order functions to specialize things further.

Arithmetic primitives expect and return integers. As in C, we reuse the code that projects two arguments to integers, but as above, we do it using higher-order functions.

```
219c <evaluation 217b>+≡ (224) <219b 220b>
arithOp: (int * int -> int) -> (value list -> value)

fun arithOp f = binaryOp (fn (NUM n1, NUM n2) => NUM (f (n1, n2))
  | _ => raise RuntimeError "integers expected")
```

We use the chunk $\langle \text{primitives} :: 220a \rangle$ to cons up all the primitives and their names into one giant list. We use that list to build the initial environment for the read-eval-print loop; it plays the same role as the unspeakable `xx` macros in file `prim.h` in the C implementation (chunk 136c).

```
220a  <primitives :: 220a>≡ (223b) 220c>
      ("+", arithOp op + ) ::
      ("-", arithOp op - ) ::
      ("*", arithOp op * ) ::
      ("/", arithOp op div) ::
```

The ML keyword `op` makes it possible to use an infix identifier as an ordinary value, so `arithOp op +` passes the value `+` (a binary function) to the function `arithOp`.

We have two kinds of predicate: an ordinary predicate takes one argument and a comparison takes two. Some comparisons apply only to integers. We reuse `embedPredicate` for the definitions.

```
220b  <evaluation 217b>+≡ (224) <219c 221b>
      predOp      : (value      -> bool) -> (value list -> value)
      comparison  : (value * value -> bool) -> (value list -> value)
      intcompare  : (int      * int  -> bool) -> (value list -> value)

      fun predOp f      = unaryOp (embedPredicate f)
      fun comparison f = binaryOp (embedPredicate f)
      fun intcompare f = comparison (fn (NUM n1, NUM n2) => f (n1, n2)
                                      | _ => raise RuntimeError "integers expected")
```

Here come the predicates. As required by the semantics of μ Scheme, equality comparison succeeds only on symbols, numbers, Booleans, and the empty list.

```
220c  <primitives :: 220a>+≡ (223b) <220a 220d>
      (<"", intcompare op <> ::
      ">", intcompare op >> ::
      "=", comparison (fn (NIL,      NIL      ) => true
                          | (NUM  n1, NUM  n2) => n1 = n2
                          | (SYM  v1, SYM  v2) => v1 = v2
                          | (BOOL b1, BOOL b2) => b1 = b2
                          | _              => false)) ::
      ("null?",  predOp (fn (NIL  ) => true | _ => false)) ::
      ("boolean?", predOp (fn (BOOL _) => true | _ => false)) ::
      ("number?", predOp (fn (NUM  _) => true | _ => false)) ::
      ("symbol?", predOp (fn (SYM  _) => true | _ => false)) ::
      ("pair?",   predOp (fn (PAIR _) => true | _ => false)) ::
      ("procedure?",
        predOp (fn (PRIMITIVE _) => true | (CLOSURE _) => true | _ => false)) ::
```

The list primitives are easy:

```
220d  <primitives :: 220a>+≡ (223b) <220c 221a>
      ("cons", binaryOp (fn (a, b) => PAIR (a, b))) ::
      ("car",  unaryOp (fn (PAIR (car, _)) => car
                          | v => raise RuntimeError
                              ("car applied to non-list " ^ valueString v))) ::
      ("cdr",  unaryOp (fn (PAIR (_, cdr)) => cdr
                          | v => raise RuntimeError
                              ("cdr applied to non-list " ^ valueString v))) ::
```


Finally, the only remaining primitives are `print` and `error`:

```
221a <primitives :: 220a>+≡ (223b) <220d>
    ("print", unaryOp (fn v => (print (valueString v^"\n"); v))) ::
    ("error", unaryOp (fn v => raise RuntimeError (valueString v))) ::
```

5.5 Evaluating definitions

As in Chapter 3, the implementation of the rules for definitions is straightforward. Function `evaldef` takes a definition and an environment, and it returns a new environment. It also takes `echo`, which is not a flag but a function used to print responses. When we see a `val`, we add the name to the environment, evaluate the right-hand side, and print the result of `showVal`, which is either the value of the right-hand side or, in the case of a binding to a function, the name of the function. As usual, `define` is syntactic sugar. When we get an expression, we print its value and bind the result to it.

```
221b <evaluation 217b>+≡ (224) <220b 222b>
```

```
evaldef : def * value ref env * (string->unit) -> value ref env
addName : name * value ref env -> value ref env
showVal : name -> exp -> value -> string
```

```
fun evaldef (d, rho, echo) =
  let <definitions of addName and showVal 221c>
  in case d
    of VAL (name, e) => let val rho = addName (name, rho)
                        val v = eval (e, rho)
                        in ( find (name, rho) := v
                          ; echo (showVal name e v)
                          ; rho
                        )
                        end
    | EXP e => let val v = eval (e, rho)
                val rho = addName ("it", rho)
                in ( find ("it", rho) := v
                  ; echo (valueString v)
                  ; rho
                )
                end
    | DEFINE (name, lambda) => evaldef (VAL (name, LAMBDA lambda), rho, echo)
    | USE filename => use readEvalPrint schemeSyntax filename rho
  end
```

```
bind      214
DEFINE    216a
eval      217b
EXP       216a
find      214
LAMBDA    215
NIL       215
NotFound  214
readEvalPrint 222b
RuntimeError
          215
schemeSyntax
          674d
unaryOp   219b
USE       216a
use       222a
VAL       216a
valueString 217a
```

The differences between `VAL` and `EXP` are subtle: for `VAL`, the rules of μ Scheme require that we add the name to environment `rho` *before* evaluating expression `e`. For `EXP`, we don't bind the name `it` until *after* evaluating the first top-level expression. Also, the results of the two kinds of bindings are printed differently.

The auxiliary functions `addName` and `showVal` are simple:

```
221c <definitions of addName and showVal 221c>≡ (221b)
    fun addName (name, rho) = (find (name, rho); rho)
                                handle NotFound _ => bind (name, ref NIL, rho)

    fun showVal name (LAMBDA _) _ = name
      | showVal name _      v = valueString v
```

The implementation of `use` is parameterized by `readEvalPrint` and `syntax` so we can share it with other interpreters. Function `use` creates a reader that does not prompt, but it uses `writeln` to be sure that responses are printed.

```
222a  <implementation of use 222a>≡ (224)
      fun use readEvalPrint syntax filename rho =
        let val fd = TextIO.openIn filename
          val defs = reader syntax noPrompts (filename, streamOfLines fd)
          fun writeln s = app print [s, "\n"]
          fun errorln s = TextIO.output (TextIO.stdErr, s ^ "\n")
        in readEvalPrint (defs, writeln, errorln) rho
        before TextIO.closeIn fd
        end
```

Functions `reader` and `streamOfLines` are defined in Appendix D. They are based on an abstraction called *streams*. A stream is like a list, except that when client code first looks at an element of a stream, the stream abstraction may do some input or output. Function `streamOfLines` produces a stream containing the lines of source code found in the named file. Function `reader syntax noPrompts` converts that stream into a stream of definitions.

5.6 The read-eval-print loop

The read-eval-print loop is built around the inner function `processDef`, which takes a definition and an environment and produces a new environment. The looping action comes from function `streamFold`, which applies `processDef` to every element of a stream of definitions, working from first to last. Function `streamFold` is the stream analog of the list function `foldl`.

As in `evaldef`, `echo` is not a flag but a function, which takes a string and may print it or do nothing. The argument `errmsg` is similar to `echo`, but it is used to issue error messages. Functions `evaldef` and `readEvalPrint` are mutually recursive, so we use `and` instead of `fun`.

```
222b  <evaluation 217b>+≡ (224) <221b>
readEvalPrint : def stream * (string->unit) * (string->unit) -> value ref env -> value ref env
      and readEvalPrint (defs, echo, errmsg) rho =
        let fun processDef (def, rho) =
          let fun continue msg = (errmsg msg; rho)
            in evaldef (def, rho, echo)
              handle IO.IOException {name, ...} => continue ("I/O error: " ^ name)
              <more read-eval-print handlers 222c>
          end
        in streamFold processDef rho defs
        end
```

```
evaldef 221b
noPrompts 668d
reader 669
streamFold 650a
streamOfLines 648b
```

The execution of `readdef` is protected by a large collection of exception handlers, each of which calls `continue` to print an error message and to return the (unchanged) environment `rho`. We reuse the same exception handlers in later interpreters, sometimes with slightly different implementations of `continue`.

The next handlers deal with problems that arise during I/O, lexical analysis, and parsing.

```
222c  <more read-eval-print handlers 222c>≡ (222b) 223a>
```

The exception `IO.IOException` is part of the Standard Basis Library.

The remaining handlers deal with problems that arise during evaluation.

```
223a  (more read-eval-print handlers 222c)+≡ (222b) <222c
      | Div          => continue "Division by zero"
      | Overflow     => continue "Arithmetic overflow"
      | RuntimeError msg => continue ("run-time error: " ^ msg)
      | NotFound n    => continue ("variable " ^ n ^ " not found")
```

Exceptions Div and Overflow are predefined.

5.7 Initializing and running the interpreter

To make a working interpreter, we need an initial environment. We create the environment by starting with the empty environment, binding the primitive operators, then adding the initial basis.

```
223b  (initialization 223b)≡ (224) 223c>
      fun initialEnv () =
        let val rho =
            foldl (fn ((name, prim), rho) => bind (name, ref (PRIMITIVE prim), rho))
                emptyEnv (<primitives :: 220a> nil)
          val basis = <ML representation of initial basis (automatically generated)>
          val defs = reader schemeSyntax noPrompts ("initial basis", streamOfList basis)
        in readEvalPrint (defs, fn _ => (), fn _ => ()) rho
        end
```

The function runInterpreter takes one argument, which tells it whether to prompt. It reads from standard input, again using streamOfLines.

```
223c  (initialization 223b)+≡ (224) <223b
      fun runInterpreter noisy =
        let val rho = initialEnv ()
          fun writeln s = app print [s, "\n"]
          fun errorln s = TextIO.output (TextIO.stdErr, s ^ "\n")
          val prompts = if noisy then stdPrompts else noPrompts
          val defs =
            reader schemeSyntax prompts ("standard input", streamOfLines TextIO.stdIn)
        in ignore (readEvalPrint (defs, writeln, errorln) rho)
        end
```

5.8 Building and exporting a program

The final step in implementing the interpreter is to create a function that looks at the command line and calls runInterpreter. With a compiler like Moscow ML or MLton, the module isn't *elaborated* until run time, so we can simply insert an irrelevant val binding, the elaboration of which has the side effect of calling main. CommandLine.arguments () returns an argument list.

```
223d  (command line 223d)≡ (224)
      fun main ["-q"] = runInterpreter false
        | main []      = runInterpreter true
        | main _       =
            TextIO.output (TextIO.stdErr, "Usage: " ^ CommandLine.name () ^ " [-q]\n")
      val _ = main (CommandLine.arguments ())
```

bind 214
continue, 214
in Typed
Impcore 246b
in Typed
μScheme 273b
in μML 330b
in μProlog 557b
in μScheme 222b
in μSmalltalk 494a
emptyEnv 214
noPrompts 668d
NotFound 214
PRIMITIVE 215
reader 669
readEvalPrint 222b
runInterpreter, 247c
in Typed
Impcore 274c
in Typed
μScheme 331b
in μML 496b
in μSmalltalk
RuntimeError, 244d
in Typed
Impcore 268a
in Typed
μScheme 287c
in μML 558a
in μProlog 215
in μScheme 469
schemeSyntax 674d
stdPrompts 668d
streamOfLines 648b
streamOfList 647c

We build the full interpreter by concatenating the parts in this order:

```
224  <mlscheme.sml 224>≡ (0—1)
      <environments 214>
      <lexical analysis 671a>
      <abstract syntax and values 215>
      <values 216b>
      <parsing 673a>
      <implementation of use 222a>
      <evaluation 217b>
      <initialization 223b>
      <command line 223d>
```

5.9 Free and bound variables

When an expression e refers to a name y that is introduced outside of e proper, we say that y is *free* in e . We often refer to the set of such names as “free variables,” even though “free name” would be more accurate. A variable in e that is introduced within e is a *bound* variable. For example, in the expression

```
(lambda (n) (+ 1 n))
```

the name $+$ is a free variable, but n is a bound variable. Every variable that appears in an expression is either free or bound.

Each variable that appears in a definition is also free or bound. For example, in

```
(define map (f xs)
  (if (null? xs)
      '()
      (cons (f (car xs)) (map f (cdr xs)))))
```

the names `null?`, `cons`, `car`, and `cdr` are free, and the names `map`, `f`, and `xs` are bound.

Free variables play a key role in implementing closures efficiently. The operational semantics for μ Scheme say that evaluating a `lambda` expression captures the *entire* environment ρ_c :

$$\frac{x_1, \dots, x_n \text{ all distinct}}{\langle \text{LAMBDA}(\langle x_1, \dots, x_n \rangle, e), \rho_c, \sigma \rangle \Downarrow \langle \langle \text{LAMBDA}(\langle x_1, \dots, x_n \rangle, e), \rho_c \rangle, \sigma \rangle} \quad (\text{MKCLOSURE})$$

Do we really need *all* the information in ρ_c ? If we look at the application rule to see how ρ_c is used, we can guess maybe not:

$$\frac{\begin{array}{c} \ell_1, \dots, \ell_n \notin \text{dom } \sigma_n \text{ (and all distinct)} \\ \langle e, \rho, \sigma \rangle \Downarrow \langle \langle \text{LAMBDA}(\langle x_1, \dots, x_n \rangle, e_c), \rho_c \rangle, \sigma_0 \rangle \\ \langle e_1, \rho, \sigma_0 \rangle \Downarrow \langle v_1, \sigma_1 \rangle \\ \vdots \\ \langle e_n, \rho, \sigma_{n-1} \rangle \Downarrow \langle v_n, \sigma_n \rangle \end{array}}{\frac{\langle e_c, \rho_c \{x_1 \mapsto \ell_1, \dots, x_n \mapsto \ell_n\}, \sigma_n \{ \ell_1 \mapsto v_1, \dots, \ell_n \mapsto v_n \} \rangle \Downarrow \langle v, \sigma' \rangle}{\langle \text{APPLY}(e, e_1, \dots, e_n), \rho, \sigma \rangle \Downarrow \langle v, \sigma' \rangle}} \quad (\text{APPLYCLOSURE})$$

How is ρ_c used? Only to evaluate the body of the `LAMBDA`. It turns out that the `MKCLOSURE` rule need not store all of ρ_c —it is enough to store only those bindings in ρ_c which refer to variables that are free in the `LAMBDA` expression. Exercise 1 asks you to prove this fact, and Exercise 2 asks you to use it to make the interpreter faster. To do these exercises, you need a precise definition of what a free variable is. We can write such a definition using a proof system.

We'll write a proof system just for expressions; the judgment form is $\boxed{y \in \text{fv}(e)}$. The notation $\text{fv}(e)$ refers to the set of all variables that appear free in e , but we'd rather not construct such a set, so we'll pronounce the judgment $y \in \text{fv}(e)$ as “ y appears free in e .” To create the proof system, we consider each syntactic form.

A literal expression has no free variables. Formally speaking, no judgment of the form $y \in \text{fv}(\text{LITERAL } v)$ can ever be proved, and we express that fact by not having a rule for literals.

An expression consisting of a single variable x has just one free variable, which is x itself:

$$\frac{}{x \in \text{fv}(\text{VAR}(x))}$$

The free variables of a SET expression include the free variables being assigned to, plus all the free variables of the right-hand side. So for a SET expression, we have two proof rules:

$$\frac{}{x \in \text{fv}(\text{SET}(x, e))} \quad \frac{y \in \text{fv}(e)}{y \in \text{fv}(\text{SET}(x, e))}$$

A variable is free in an IF expression if and only if it is free in one of the subexpressions:

$$\frac{y \in \text{fv}(e_1)}{y \in \text{fv}(\text{IF}(e_1, e_2, e_3))} \quad \frac{y \in \text{fv}(e_2)}{y \in \text{fv}(\text{IF}(e_1, e_2, e_3))} \quad \frac{y \in \text{fv}(e_3)}{y \in \text{fv}(\text{IF}(e_1, e_2, e_3))}$$

A variable is also free in a WHILE expression if and only if it is free in one of the subexpressions:

$$\frac{y \in \text{fv}(e_1)}{y \in \text{fv}(\text{WHILE}(e_1, e_2))} \quad \frac{y \in \text{fv}(e_2)}{y \in \text{fv}(\text{WHILE}(e_1, e_2))}$$

And the same for BEGIN:

$$\frac{y \in \text{fv}(e_i)}{y \in \text{fv}(\text{BEGIN}(e_1, \dots, e_n))}$$

A variable is free in an application if and only if it is free in the function or in one of the arguments:

$$\frac{y \in \text{fv}(e)}{y \in \text{fv}(\text{APPLY}(e, e_1, \dots, e_n))} \quad \frac{y \in \text{fv}(e_i)}{y \in \text{fv}(\text{APPLY}(e, e_1, \dots, e_n))}$$

Finally, we get to an interesting case! A variable is free in a LAMBDA expression provided that it is free in the body, and it is *not* one of the arguments:

$$\frac{y \in \text{fv}(e) \quad y \notin \{x_1, \dots, x_n\}}{y \in \text{fv}(\text{LAMBDA}(\langle x_1, \dots, x_n \rangle, e))}$$

The rules for the various LET forms require care. The free variables of an ordinary LET are the free variables of the right-hand sides of all the bindings, plus those free variables of the body which are not LET-bound.

$$\frac{y \in \text{fv}(e_i)}{y \in \text{fv}(\text{LET}(\langle x_1, e_1, \dots, x_n, e_n \rangle, e))} \quad \frac{y \in \text{fv}(e) \quad y \notin \{x_1, \dots, x_n\}}{y \in \text{fv}(\text{LET}(\langle x_1, e_1, \dots, x_n, e_n \rangle, e))}$$

The similarity between the second LET rule and the LAMBDA rule shows a kinship between LET and LAMBDA.

The rule for LETREC is very similar, except that in a LETREC, the bound names x_i are never free:

$$\frac{y \in \text{fv}(e_i) \quad y \notin \{x_1, \dots, x_n\}}{y \in \text{fv}(\text{LET}(\langle x_1, e_1, \dots, x_n, e_n \rangle, e))} \quad \frac{y \in \text{fv}(e) \quad y \notin \{x_1, \dots, x_n\}}{y \in \text{fv}(\text{LET}(\langle x_1, e_1, \dots, x_n, e_n \rangle, e))}$$

As is typical, it is a nuisance to try to write the rule for LETSTAR directly. Instead, we treat a LETSTAR expression as a set of nested LET expressions, each containing just one binding. And an empty LETSTAR behaves just like its body.

$$\frac{y \in \text{fv}(\text{LET}(\langle x_1, e_1 \rangle, \text{LETSTAR}(\langle x_2, e_2, \dots, x_n, e_n \rangle, e)))}{y \in \text{fv}(\text{LETSTAR}(\langle x_1, e_1, \dots, x_n, e_n \rangle, e))} \quad \frac{y \in \text{fv}(e)}{y \in \text{fv}(\text{LETSTAR}(\langle \rangle, e))}$$

5.10 Exercises

1. In this exercise, you prove that the evaluation of an expression doesn't depend on arbitrary bindings in the environment, but only on the bindings of the expression's free variables.

If X is a set of variables, we can ask what happens to an environment ρ if we remove the bindings of all the names that are *not* in the set X . The modified environment is written $\rho|_X$, and it is called the *restriction* of ρ to X . The exercise is to prove, by structural induction on derivations, that if $\langle e, \rho, \sigma \rangle \Downarrow \langle v, \sigma \rangle$, then $\langle e, \rho|_{\text{fv}(e)}, \sigma \rangle \Downarrow \langle v, \sigma \rangle$. To structure the proof, I recommend you introduce a definition and a lemma.

- Define $\rho \sqsubseteq \rho'$ to mean that $\text{dom } \rho \subseteq \text{dom } \rho'$, and $\forall x \in \text{dom } \rho : \rho(x) = \rho'(x)$. The domain of ρ' contains the domain of ρ , and on their common domain, they agree. We might say that ρ' *extends* ρ .
- If $X \subseteq X' \subseteq \text{dom } \rho$, then $\rho|_X \sqsubseteq \rho|_{X'} \sqsubseteq \rho$.

These tools are useful, because except for LET forms and LAMBDA expressions, if an expression e has a subexpression e_i , then $\text{fv}(e_i) \subseteq \text{fv}(e)$.

A reasonable induction hypothesis for the proof might be that if that if $\langle e, \rho, \sigma \rangle \Downarrow \langle v, \sigma \rangle$, and if $\rho|_{\text{fv}(e)} \sqsubseteq \rho'$, then $\langle e, \rho', \sigma \rangle \Downarrow \langle v, \sigma \rangle$.

2. The payoff for the proof in Exercise 1 is that we can use it to optimize code. In chunk 218b, a LAMBDA expression is evaluated by capturing a full environment ρ . Modify the code to capture a restricted environment that contains only the free variables of the LAMBDA expression. That is, instead of allocating the closure $\langle \text{LAMBDA}(\langle x_1, \dots, x_n \rangle, e), \rho \rangle$, allocate $\langle \text{LAMBDA}(\langle x_1, \dots, x_n \rangle, e), \rho|_X \rangle$, where $X = \text{fv}(\text{LAMBDA}(\langle x_1, \dots, x_n \rangle, e))$.
3. Change the evaluation of lambda expressions so that the result is a function of type `value list -> value`. On the strength of this trick, rename PRIMITIVE to PROCEDURE, and eliminate CLOSURE from the interpreter.
What are the drawbacks to this simplification?
4. Revisit the material on proofs and derivations in Section 2.6.
 - (a) Devise a representation, in Standard ML, of the judgments of the semantics for μScheme .
 - (b) Devise a representation, in Standard ML, of derivations that use the operational semantics of μScheme .
 - (c) Change the `eval` function of the μScheme interpreter to return a *derivation* instead of a value.
5. Using the representation of derivations in Exercise 4b, write a *proof checker* that tells whether a given tree represents a valid derivation.

