

Control flow graphs and loop optimizations

Friday, November 7, 14

Agenda

- Building control flow graphs
- Low level loop optimizations
 - Code motion
 - Strength reduction
 - Unrolling
- High level loop optimizations
 - Loop fusion
 - Loop interchange
 - Loop tiling

Friday, November 7, 14

Moving beyond basic blocks

- Up until now, we have focused on single basic blocks
- What do we do if we want to consider larger units of computation
 - Whole procedures?
 - Whole program?
- Idea: capture *control flow* of a program
- How control transfers between basic blocks due to:
 - Conditionals
 - Loops

Friday, November 7, 14

Representation

- Use standard three-address code
- Jump targets are labeled
- Also label beginning/end of functions
- Want to keep track of *targets of jump statements*
 - Any statement whose execution may immediately follow execution of jump statement
 - *Explicit* targets: targets mentioned in jump statement
 - *Implicit* targets: statements that follow conditional jump statements
 - The statement that gets executed if the branch is not taken

Friday, November 7, 14

Running example

```
A = 4
t1 = A * B
repeat {
  t2 = t1 / C
  if (t2 ≥ W) {
    M = t1 * k
    t3 = M + I
  }
  H = I
  M = t3 - H
} until (T3 ≥ 0)
```

Friday, November 7, 14

Running example

```
1      A = 4
2      t1 = A * B
3  L1:  t2 = t1 / C
4      if t2 < W goto L2
5      M = t1 * k
6      t3 = M + I
7  L2:  H = I
8      M = t3 - H
9      if t3 ≥ 0 goto L3
10     goto L1
11  L3:  halt
```

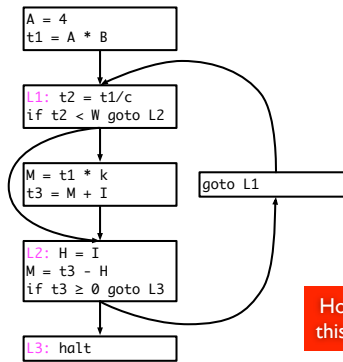
Friday, November 7, 14

Control flow graphs

- Divides statements into *basic blocks*
- Basic block: a maximal sequence of statements $l_0, l_1, l_2, \dots, l_n$ such that if l_i and l_{i+1} are two adjacent statements in this sequence, then
 - The execution of l_i is always immediately followed by the execution of l_{i+1}
 - The execution of l_{i+1} is always immediate preceded by the execution of l_i
- Edges between basic blocks represent potential flow of control

Friday, November 7, 14

CFG for running example



How do we build this automatically?

Friday, November 7, 14

Constructing a CFG

- To construct a CFG where each node is a basic block
 - Identify *leaders*: first statement of a basic block
 - In program order, construct a block by appending subsequent statements up to, but not including, the next leader
- Identifying leaders
 - First statement in the program
 - Explicit target of any conditional or unconditional branch
 - Implicit target of any branch

Friday, November 7, 14

Partitioning algorithm

- Input: set of statements, $stat(i) = i^{th}$ statement in input
- Output: set of *leaders*, set of basic blocks where $block(x)$ is the set of statements in the block with leader x
- Algorithm


```

leaders = {1} //Leaders always includes first statement
for i = 1 to |n| //|n| = number of statements
    if stat(i) is a branch, then
        leaders = leaders ∪ all potential targets
end for
worklist = leaders
while worklist not empty do
    x = remove earliest statement in worklist
    block(x) = {x}
    for (i = x + 1; i ≤ |n| and i ∉ leaders; i++)
        block(x) = block(x) ∪ {i}
    end for
end while
            
```

Friday, November 7, 14

Running example

```

1      A = 4
2      t1 = A * B
3  L1:  t2 = t1 / C
4      if t2 < W goto L2
5      M = t1 * k
6      t3 = M + I
7  L2:  H = I
8      M = t3 - H
9      if t3 ≥ 0 goto L3
10     goto L1
11  L3:  halt
            
```

Leaders =
Basic blocks =

Friday, November 7, 14

Running example

```

1      A = 4
2      t1 = A * B
3  L1:  t2 = t1 / C
4      if t2 < W goto L2
5      M = t1 * k
6      t3 = M + I
7  L2:  H = I
8      M = t3 - H
9      if t3 ≥ 0 goto L3
10     goto L1
11  L3:  halt
            
```

Leaders = {1, 3, 5, 7, 10, 11}
Basic blocks = {{1, 2}, {3, 4}, {5, 6}, {7, 8, 9}, {10}, {11}}

Friday, November 7, 14

Putting edges in CFG


- There is a directed edge from B_1 to B_2 if
 - There is a branch from the last statement of B_1 to the first statement (leader) of B_2
 - B_2 immediately follows B_1 in program order and B_1 does not end with an unconditional branch
- Input: *block*, a sequence of basic blocks
- Output: The CFG


```

for i = 1 to |block|
  x = last statement of block(i)
  if stat(x) is a branch, then
    for each explicit target y of stat(x)
      create edge from block i to block y
    end for
  if stat(x) is not unconditional then
    create edge from block i to block i+1
  end for
      
```

Friday, November 7, 14

Result




Friday, November 7, 14

Discussion

- Some times we will also consider the *statement-level* CFG, where each node is a statement rather than a basic block
- Either kind of graph is referred to as a CFG
- In statement-level CFG, we often use a node to explicitly represent *merging* of control
- Control merges when two different CFG nodes point to the same node
- Note: if input language is *structured*, front-end can generate basic block directly
- “GOTO considered harmful”

Friday, November 7, 14

Statement level CFG



Friday, November 7, 14

Loop optimization

- Low level optimization
 - Moving code around in a single loop
 - Examples: loop invariant code motion, strength reduction, loop unrolling
- High level optimization
 - Restructuring loops, often affects multiple loops
 - Examples: loop fusion, loop interchange, loop tiling

Friday, November 7, 14

Low level loop optimizations

- Affect a single loop
- Usually performed at three-address code stage or later in compiler
- First problem: identifying loops
 - Low level representation doesn't have loop statements!

Friday, November 7, 14


Identifying loops

- First, we must identify **dominators**
 - Node **a** dominates node **b** if every possible execution path that gets to **b** must pass through **a**
- Many different algorithms to calculate dominators – we will not cover how this is calculated
- A **back edge** is an edge from **b** to **a** when **a** dominates **b**
- The target of a back edge is a **loop header**

Friday, November 7, 14

Natural loops

- Will focus on **natural loops** – loops that arise in structured programs
- For a node **n** to be in a loop with header **h**
 - n** must be dominated by **h**
 - There must be a path in the CFG from **n** to **h** through a back-edge to **h**
- What are the back edges in the example to the right? The loop headers? The natural loops?



Friday, November 7, 14

Loop invariant code motion

- Idea: some expressions evaluated in a loop never change; they are **loop invariant**
- Can move loop invariant expressions outside the loop, store result in temporary and just use the temporary in each iteration
- Why is this useful?

Friday, November 7, 14

Identifying loop invariant code

- To determine if a statement **s: a = b op c** is loop invariant, find all definitions of **b** and **c** that reach **s**
 - A statement **t** defining **b** reaches **s** if there is a path from **t** to **s** where **b** is not re-defined
- s** is loop invariant if both **b** and **c** satisfy one of the following
 - it is constant
 - all definitions that reach it are from outside the loop
 - only one definition reaches it and that definition is also loop invariant

Friday, November 7, 14

Moving loop invariant code

- Just because code is loop invariant doesn't mean we can move it!


```

do
  if (*)
    break
a = 5
while (*)
  c = a;
      
```

```

for (...)
  if (*)
    a = 5
  else
    a = 6
      
```

```

a = 5;
for (...)
  if (*)
    a = 4 + c
  b = a
      
```
- We can move a loop invariant statement **a = b op c** if
 - The statement dominates all loop exits where **a** is live
 - There is only one definition of **a** in the loop
 - a** is not live before the loop
- Move instruction to a **preheader**, a new block put right before loop header

Friday, November 7, 14

Strength reduction

- Like strength reduction peephole optimization
- Peephole: replace expensive instruction like **a * 2** with **a << 1**
- Replace expensive instruction, multiply, with a cheap one, addition
 - Applies to uses of an **induction variable**
 - Opportunity: array indexing

```

for (i = 0; i < 100; i++)
  A[i] = 0;
  
```

↓

```

i = 0;
L2: if (i >= 100) goto L1
  j = 4 * i + &A
  *j = 0;
  i = i + 1;
  goto L2
L1:
  
```

Friday, November 7, 14

Strength reduction

- Like strength reduction peephole optimization
- Peephole: replace expensive instruction like $a * 2$ with $a << 1$
- Replace expensive instruction, multiply, with a cheap one, addition
- Applies to uses of an *induction variable*
- Opportunity: array indexing

```
for (i = 0; i < 100; i++)
    A[i] = 0;

↓

i = 0; k = &A;
L2: if (i >= 100) goto L1
    j = k;
    *j = 0;
    i = i + 1; k = k + 4;
    goto L2
L1:
```

Friday, November 7, 14

Induction variables

- A *basic induction variable* is a variable j
 - whose only definition within the loop is an assignment of the form $j = j \pm c$, where c is loop invariant
 - Intuition: the variable which determines number of iterations is usually an induction variable
- A *mutual induction variable* i may be
 - defined once within the loop, and its value is a linear function of some other induction variable j such that

$$i = c_1 * j \pm c_2 \text{ or } i = j / c_1 \pm c_2$$
 where c_1, c_2 are loop invariant
- A *family* of induction variables include a basic induction variable and any related mutual induction variables

Friday, November 7, 14

Strength reduction algorithm

- Let i be an induction variable in the family of the basic induction variable j , such that $i = c_1 * j + c_2$
 - Create a new variable i'
 - Initialize in preheader

$$i' = c_1 * j + c_2$$
 - Track value of j . After $j = j + c_3$, perform

$$i' = i' + (c_1 * c_3)$$
 - Replace definition of i with

$$i = i'$$
- Key: c_1, c_2, c_3 are all loop invariant (or constant), so computations like $(c_1 * c_3)$ can be moved outside loop

Friday, November 7, 14

Linear test replacement

- After strength reduction, the loop test may be the only use of the basic induction variable
- Can now eliminate induction variable altogether
- Algorithm
 - If only use of an induction variable is the loop test and its increment, and if the test is always computed
 - Can replace the test with an equivalent one using one of the mutual induction variables

```
i = 2
for (; i < k; i++)
    j = 50 * i
    ... = j

↓ Strength reduction

i = 2; j' = 50 * i
for (; i < k; i++, j' += 50)
    ... = j'

↓ Linear test replacement

i = 2; j' = 50 * i
for (; j' < 50 * k; j' += 50)
    ... = j'
```

Friday, November 7, 14

Loop unrolling

- Modifying induction variable in each iteration can be expensive
- Can instead *unroll* loops and perform multiple iterations for each increment of the induction variable
- What are the advantages and disadvantages?

```
for (i = 0; i < N; i++)
    A[i] = ...

↓ Unroll by factor of 4

for (i = 0; i < N; i += 4)
    A[i] = ...
    A[i+1] = ...
    A[i+2] = ...
    A[i+3] = ...
```

Friday, November 7, 14


High level loop optimizations

- Many useful compiler optimizations require *restructuring* loops or sets of loops
 - Combining two loops together (*loop fusion*)
 - Switching the order of a nested loop (*loop interchange*)
 - Completely changing the traversal order of a loop (*loop tiling*)
- These sorts of high level loop optimizations usually take place at the AST level (where loop structure is obvious)

Friday, November 7, 14

Cache behavior

- Most loop transformations target cache performance
- Attempt to increase *spatial* or *temporal* locality
- Locality can be exploited when there is reuse of data (for temporal locality) or recent access of nearby data (for spatial locality)
- Loops are a good opportunity for this: many loops iterate through matrices or arrays
- Consider matrix-vector multiply example
- Multiple traversals of vector: opportunity for spatial and temporal locality
- Regular access to array: opportunity for spatial locality



$$y = Ax$$

```
for (i = 0; i < N; i++)
  for (j = 0; j < N; j++)
    y[i] += A[i][j] * x[j]
```

Friday, November 7, 14


Loop fusion

```
do l = 1, n
  c[l] = a[l]
end do
do l = 1, n
  b[l] = a[l]
end do
```



- Combine two loops together into a single loop
- Why is this useful?
- Is this always legal?


```
do l = 1, n
  c[l] = a[l]
  b[l] = a[l]
end do
```



Friday, November 7, 14

Loop interchange

- Change the order of a nested loop
- This is not always legal – it changes the order that elements are accessed!
- Why is this useful?
- Consider matrix-matrix multiply when A is stored in column-major order (i.e., each column is stored in contiguous memory)




```
for (i = 0; i < N; i++)
  for (j = 0; j < N; j++)
    y[i] += A[i][j] * x[j]
```

Friday, November 7, 14

Loop interchange

- Change the order of a nested loop
- This is not always legal – it changes the order that elements are accessed!
- Why is this useful?
- Consider matrix-matrix multiply when A is stored in column-major order (i.e., each column is stored in contiguous memory)



```
for (j = 0; j < N; j++)
  for (i = 0; i < N; i++)
    y[i] += A[i][j] * x[j]
```


Friday, November 7, 14

Loop tiling

- Also called “loop blocking”
- One of the more complex loop transformations
- Goal: break loop up into smaller pieces to get spatial and temporal locality
- Create new inner loops so that data accessed in inner loops fit in cache
- Also changes iteration order, so may not be legal

```
for (i = 0; i < N; i++)
  for (j = 0; j < N; j++)
    y[i] += A[i][j] * x[j]

for (ii = 0; ii < N; ii += B)
  for (jj = 0; jj < N; jj += B)
    for (i = ii; i < ii+B; i++)
      for (j = jj; j < jj+B; j++)
        y[i] += A[i][j] * x[j]
```




Friday, November 7, 14

Loop tiling

- Also called “loop blocking”
- One of the more complex loop transformations
- Goal: break loop up into smaller pieces to get spatial and temporal locality
- Create new inner loops so that data accessed in inner loops fit in cache
- Also changes iteration order, so may not be legal


```
for (i = 0; i < N; i++)
  for (j = 0; j < N; j++)
    y[i] += A[i][j] * x[j]

for (ii = 0; ii < N; ii += B)
  for (jj = 0; jj < N; jj += B)
    for (i = ii; i < ii+B; i++)
      for (j = jj; j < jj+B; j++)
        y[i] += A[i][j] * x[j]
```



Friday, November 7, 14

In a real (Itanium) compiler



Friday, November 7, 14

Loop transformations

- Loop transformations can have dramatic effects on performance
- Doing this legally and automatically is very difficult!
- Researchers have developed techniques to determine legality of loop transformations and automatically transform the loop
 - Techniques like *unimodular transform framework* and *polyhedral framework*
- These approaches will get covered in more detail in advanced compilers course

Friday, November 7, 14