


# ECE595: Introduction to Operating Systems

Y. Charlie Hu

1/9/2017




## Missing piece of the puzzle



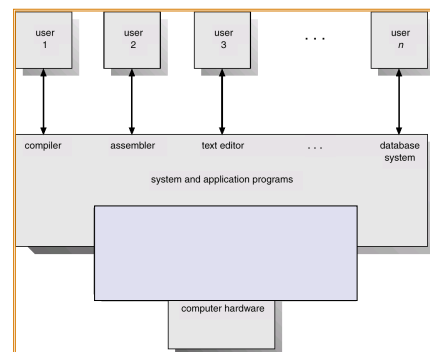
2

## A Typical Computer from a Hardware Point of View



3


## Computer System Components



4

## Example: programming hard drive

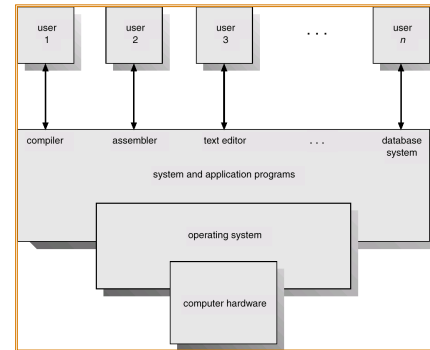
- Physical reality
  - Block oriented (e.g. 512 bytes)
  - Physical sector numbers
  - No protection among users of the system
  - Data might be corrupted if machine crashes
- Programming:
  - Loading values into special device registers



"I will save prog. assignment 1 solution on platter 5, track 8739, sector 3-4."

5

## Computer System Components



6

## Example: programming hard drive

- Physical reality
  - Block oriented
  - Physical sector numbers
  - No protection among users of the system
  - Data might be corrupted if machine crashes
- Programming:
  - Loading values into special device registers
- File system abstraction
  - Byte oriented
  - Named files
  - Users protected from each other
  - Robust to machine failures
  - Programming
    - open/read/write/close



"I will save prog. assignment 1 solution on platter 5, track 8739, sector 3-4."

"My assignment 1 solution is in ~ychu/proj1/process.c."

7

## Brief History of Computer Systems (1)

- In the beginning, 1 user/program at a time (by operator)
- Simple **batch** systems were 1<sup>st</sup> real OS
  - Spooling and buffering** allowed jobs to be read ahead of time
- Multiprogramming** systems provided increased utilization (throughput)
  - multiple runnable jobs loaded in memory
  - overlap I/O with computation
  - benefit from asynchronous I/O devices (interrupt, DMA, ...)
  - 1<sup>st</sup> instance where the OS must schedule resources
    - CPU scheduling
    - Memory management
    - Protection

8

## Brief History of Computer Systems (2)



- **Timesharing** systems support interactive use
  - Logical extension of multiprogramming
  - Permits interactive work
    - Each user feels he/she has the entire machine
  - Optimize response time by frequent time-slicing multiple jobs
- More complex than multiprogramming OS
  - In addition to CPU scheduling, memory management, protection
  - Virtual memory to allow part of the job be in memory
  - File system (needed by interactive use)
  - Job communication, synchronization
  - Handling deadlocks
- Most systems today are timesharing systems (focus of this class) <sup>9</sup>

## What is an OS?



“Code” that *sits between*:

- programs & hardware
- different programs
- different users

But what does it do/achieve?

10

## What is an OS?



- **Resource manager**
- **Extended (abstract) machine**

Makes computers *efficient* and *easy to use*

- (will have a 3<sup>rd</sup> def based on pragmatics)

11

## What is an OS?



Resource manager (answer1)

- Allocation
- Reclamation
- Protection

12

## What is an OS?

Resource manager

- Allocation
- Reclamation
- Protection

Finite resources  
Competing demands

Examples:

- CPU
- Memory
- Disk
- Network

13

## What is an OS?

Resource manager

- Allocation
- Reclamation
- Protection

“The OS giveth  
The OS taketh away”

Implied at termination  
Involuntary at run time  
Cooperative (yield cpu)

14

## What is an OS?

Resource manager

- Allocation
- Reclamation
- Protection

“You can’t hurt me  
I can’t hurt you”

Implies some degree of  
safety & security

15

## What is an OS?

Extended (abstract) machine (answer 2)

- Much more ideal environment than the hardware
  - Ease to use
  - Fair (well-behaved)
  - Supporting backward-compatibility
  - Reliable
  - Secure
- Illusion of infinite, private (reliable, secure) resources
  - Single processor → many separate processors
  - Single memory → many separate, larger memories

16

## Is there a perfect OS? (resource manager, abstract machine)

Efficiency  
Fairness

Portability  
Interfaces

Security  
Robustness

- Conflicting goals
  - Fairness vs efficiency
  - Efficiency vs portability
  - ...
- Furthermore, ...

17

## Hardware is evolving...

- 60's-70's - Mainframes
  - Rise of IBM
- 70's - 80's - Minicomputers
  - Rise of Digital Equipment
- 80's - 90's - PCs
  - Rise of Intel, Microsoft
- 90's - 00's - handheld/portable systems (laptops)
- 2007 - today -- mobile systems (smartphones), Internet of Things
  - Rise of iPhone, Android

18

## Implications on OS Design Goals: Historical Comparison

	Mainframe	Mini	Micro/ Mobile
System \$/ worker	10:1 – 100:1	10:1 – 1:1	1:10-1:100
Performance goal	System utilization	Overall cost	Worker productivity
Functionality goal	Maximize utilization	Features	Ease of Use

19

## Hardware is evolving (cont) ...

- New architectures
  - Multiprocessors
  - 32-bit vs. 64-bit
  - Multi-core

20

## May You Live in Interesting Times...



- Processor speed doubles in 18 months
    - Number of cores per chip doubles in 24 months
  - Disk capacity doubles every 12 months
  - Global bandwidth doubles every 6 months
- Performance/cost “sweet spot” constantly decaying

*\* Does human productivity ever double?*

21

## Applications are also evolving...



- New applications
  - Scientific computing
  - Computer games
  - Java
  - WWW (web servers, browsers)
  - Networked games
  - Peer-to-peer
  - Web 2.0 (search, youtube, social network, ...)
  - Mobile apps (1.5+ million iPhone, Android apps each)
  - ...

22

## Implications to OS Design



- Constant evolution of hardware and applications continuously reshape
  - OS design goals (performance vs. functionality)
  - OS design performance/cost tradeoffs
- Any magic bullet to good OS design?

23

## There is no magic in OS design



### This is Engineering

- Imperfection
- Tradeoffs (perf/func)
- Constraints
  - hardware, cost, time, power
- Optimizations

### Nothing's Permanent

- High rate of change
  - Hardware
  - Applications
- Cost / benefit analyses
- One good news:
  - Inertia of a few design principles

24

## Separating Policies from Mechanisms



A fundamental design principle in Computer Science

**Mechanism** – tool/implementation to achieve some effect

**Policy** – decisions on what effect should be achieved

Example – CPU scheduling:

- All users treated equally
- All program instances treated equally
- Preferred users treated better

Separation leads to flexibility!

25