

System Calls, Interrupt

ECE595, Jan 11

Y. Charlie Hu



1

Roadmap

- System calls
- Interrupt



2

[lec1] What is an OS?

- Resource manager
 - Manage shared resources (CPU, mem, I/O, networking)
- Extended (abstract) machine



3


Dual-Mode Operation (ECE437)

- OS manages shared resources
- OS protects programs from other programs
 - OS needs to be “privileged”
- Dual-mode operation of hardware
 - Kernel mode – can run *privileged* instructions
 - User mode – can only run *non-privileged* instructions



4

Typical Unix OS Structure



5

Dual-Mode Operation

- OS manages shared resources
- OS protects programs from other programs
- ➔ OS needs to be privileged

- If OS manages shared resources, how does a user program request for accessing shared resources (e.g. hard drive)?


6

System calls

- Interface between a *process* and the operating system kernel
 - Kernel manages shared resources & exports interface
 - Process requests for access to shared resources
- Generally available as assembly-language instructions
- Directly invoked in many languages (C, C++, Perl)
 - Who is helping out here?

7

Typical Unix OS Structure



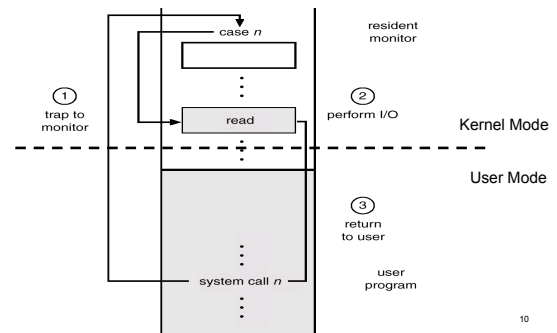
8

Example

- Given the I/O instructions are privileged, how does the user program perform I/O?
 - open()
 - read()
 - write()
 - close()

9

Use of a system call to perform I/O



10

System calls

- Categories
 - Process management
 - Memory management
 - File management
 - Device management
 - Networking


11

System calls in Linux (man syscalls)

- SYSCALLS(2) Linux Programmer's Manual SYSCALLS(2)
- NAME
 - none - list of all system calls
- SYNOPSIS
 - Linux 2.4 system calls.
- DESCRIPTION
 - The system call is the fundamental interface between an application and the Linux kernel. As of Linux 2.4.17, there are 1100 system calls listed in /usr/src/linux/include/asm-*/unistd.h. This man page lists those that are common to most platforms (providing hyperlinks if you read this with a browser).
 - _llseek(2), _newselect(2), _sysctl(2), accept(2), access(2), acct(2), adjtimex(2), afs_syscall, alarm(2), bdflush(2), bind(2), break, brk(2), cacheflush(2), capget(2), capset(2), chdir(2), chmod(2), chown(2), chown32, chroot(2), clone(2), close(2), connect(2), creat(2), create_module(2), delete_module(2), dup(2), dup2(2), execve(2), exit(2), fchdir(2), fchmod(2), fchown(2), fchown32, fcntl(2), fcntl64, fdata-
 -

12

Transition from user to kernel mode



Project 1 (week 2)

- DLXOS setup
- Trap (drop into the kernel) implementation
 - Used in later projects
 - Adding a new system call (getpid)


14

Roadmap

- System calls
- Interrupt

15

A Typical Computer from a Hardware Point of View



Concurrency & Unexpected Events



- How do humans handle unexpected events?
 - Assume
 - mail delivered to my mailbox continuously
 - I have a secretary
 - Do I have mail now? (need to be processed quickly)
- Poll vs. interrupt
- Which one is more efficient?
 - Assume 1 interrupt more costly than 1 poll
 - If I have one mail per day?
 - If I have lots of mail per delivery?

17

Interrupt (ECE437)



- A mechanism for
 - coordination between concurrently operating units of a computer system (e.g. CPU and I/O devices)
 - for responding to specific conditions within a computer
- Results in transfer of flow of control (to interrupt handler in the OS), **forced by hardware**

18

Two types of Interrupts



- **Hardware interrupts**
 - timers
 - I/O devices: keyboards, NICs, etc.
- **Software interrupts** (aka. *trap, exception*)
 - an error (floating point exception)
 - a *system call* requesting OS for special services (e.g. I/O)

19

Handling interrupts (ECE437)



- Incoming interrupts are disabled (at this and lower priority levels) while the interrupt is being processed to prevent a lost interrupt
- Interrupt architecture must save the address of the interrupted instruction
- Interrupt transfers control to the **interrupt service routine**
 - generally, through the interrupt vector, which contains the addresses of all the service routines
- If interrupt routine modifies process state (register values)
 - save the current state of the CPU (registers and the program counter) on the system stack
 - restore the state before returning
- Interrupts are re-enabled after servicing current interrupt
- Resume the interrupted instruction

20

X86 Interrupt and Exceptions (1)

Vector #	Mnemonic	Description	Type
0	#DE	Divide error (by zero)	Fault
1	#DB	Debug	Fault/trap
2		Non-Maskable interrupt	Interrupt
3	#BP	Breakpoint	Trap
4	#OF	Overflow	Trap
5	#BR	BOUND range exceeded	Trap
6	#UD	Invalid opcode	Fault
7	#NM	Device not available	Fault
8	#DF	Double fault	Abort
9		Coprocessor segment overrun	Fault
10	#TS	Invalid TSS	



21

X86 Interrupt and Exceptions (2)

Vector #	Mnemonic	Description	Type
11	#NP	Segment not present	Fault
12	#SS	Stack-segment fault	Fault
13	#GP	General protection	Fault
14	#PF	Page fault	Fault
15		Reserved	Fault
16	#MF	Floating-point error (math fault)	Fault
17	#AC	Alignment check	Fault
18	#MC	Machine check	Abort
19-31		Reserved	
32-255		User defined	Interrupt

Vector 128 is for system calls

22



“Modern OSes are interrupt driven”

- “An OS is a giant interrupt handler!” (Def 3)
 - Timer interrupt → Context switches in multiprogramming
 - (unexpected) I/O interrupts
 - System calls (traps) to switch from user to supervisor mode
- At the lowest level an OS is just a bunch of interrupt service routings
 - Each routine simply returns to whatever was executing before it was interrupted
 - A user process
 - An OS process
 - Another interrupt routine
 - Else infinite wait loop

23



Reading Assignment

- We will dive in process management next week
- We will go through Project 1 on Friday
- Reading assignment: OSC Chapters 1-2

24

