

## Process Synchronization

ECE595

Jan 20

Y. Charlie Hu



1

## Roadmap

- What is a process?
- Concurrent processes
- Process creation
- Non-preemptive CPU scheduling
- Process synchronization



2

## Multiprogramming needs CPU scheduling

- Without any hardware support, what can the OS do to a running process?



3

## System calls may trigger Scheduler

- Block – wait on some event/resource
  - Network packet arrival (e.g., `recv()`)
  - Keyboard, mouse input (e.g., `getchar()`)
  - Disk activity completion (e.g., `read()`)
- Yield – give up running for now




4

## Non-Preemptive Scheduler

- A non-preemptive scheduler: a scheduler that is only invoked by explicit block/yield calls, or terminations
  - Only method when there is no timer!
- The simplest form  
Scheduler:
  - save current process state (into PCB)
  - choose next process to run
  - dispatch (load PCB and run)
- Used in Windows 3.1, Mac OS

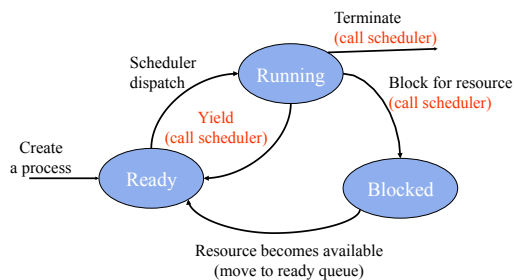
5

## Our Friend -- the Transition Diagram



6

## Process State Transition of Non-Preemptive Scheduling



How does a process terminate itself?

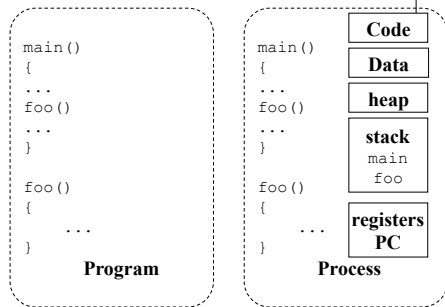
7

## Context Switch

- Definition:  
switching the CPU to run another process, which involves (1) saving the state of the old process and (2) loading the state of the new process
- What state?

8

## [lec3] Program vs. Process



9

## Context Switch

- Definition:  
switching the CPU to run another process, which involves (1) saving the state of the old process and (2) loading the state of the new process
- What state?
  - What about L1/L2 cache content?

10

## Context Switch

- Definition:  
switching the CPU to another process, which involves saving the state of the old process and load the state of the new process
- What state?
- Where to store them?

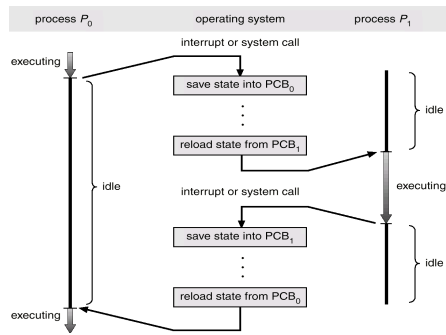
11

## [lec3] Process Control Block (Process Table)

- Process management info
  - State (ready, running, blocked)
  - PC & Registers, parents, etc
  - CPU scheduling info (priorities, etc.)
- Memory management info
  - Segments, page table, stats, etc
- I/O and file management
  - Communication ports, directories, file descriptors, etc.

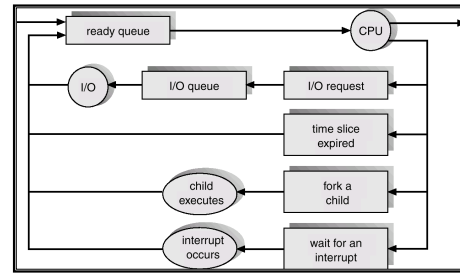
12

## Context Switch



13

## Process Scheduling Diagram



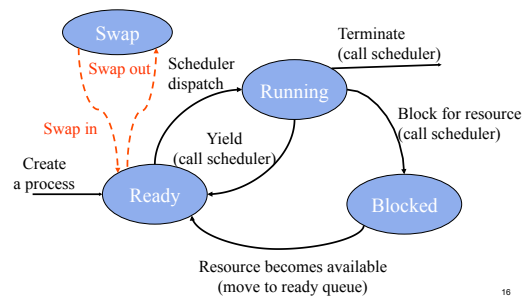
14

## Physical Memory & Multiprogramming

- Want to run many programs
- Programs need memory to run
- Memory is a scarce resource
- What happens when  $M(a) + M(b) + M(c) > \text{physical mem?}$

15

## Add Job Swapping to State Transition Diagram



16

## Process synchronization

- Cooperating processes may share data via
  - shared address space (code, data, heap) by using threads
  - shared memory objects (used in lab2)
  - Files
  - (Sending messages)
- What can happen if processes try to access shared data (address) concurrently?
  - Sharing bank account with sibling:  
At 3pm: If (balance > \$10) withdraw \$10
- How hard is the solution?

17

## “Got Milk?”

- “Too much milk” problem

18

## Process synchronization needs help from OS!

19

## ***Mutual exclusion & Critical Section***

- ***Critical section*** – a section of code, or collection of operations, in which only one process shall be executing at a given time
- ***Mutual exclusion*** - mechanisms that ensure that only one person or process is doing certain things at one time (others are excluded)

20

## Example of Critical Section

- Concurrent accesses to shared variables, at least one of which is write

P0:	P1:	P2:
Read note;	Read note;	write note;
...	...	...

21

## Desirable properties of MuEx

- *Fair*: if several processes are waiting, let each in eventually
- *Efficient*: don't use up substantial amounts of resources when waiting (e.g. no busy waiting)
- *Simple*: should be easy to use (e.g. just bracket the critical sections)

22

## Desirable properties of processes using MuEx

- Always lock before manipulating shared data
- Always unlock after manipulating shared data
- Do not lock again if already locked
- Do not unlock if not locked by you
- Do not spend large amounts of time in critical section

23

## Mutual Exclusion provided by OS (and language/compiler)

- Locks
  - Alone can solve simple problems
- More advanced
  - Semaphore
  - Lock and condition variable
    - Lock alone is not flexible enough
  - Monitor
- Each primitive itself is atomic!

24

## Lock (aka mutex)

Init: lock = 1; // 0 means held; 1 means free

```
lock_acquire(lock)      lock_release(lock)

{                        {
    while (lock==0);      if (lock == 0)
    lock--;               lock++;
}                        }
```

- Each primitive is atomic
- In reality, lock is not implemented as above!
  - The waiting process is put to sleep

25

## “Too much milk” problem with locks

```
Acquire(lock);
if (noMilk)
    buy milk;
Release(lock); } Critical Section
```

- What is the problem with this solution?

26

## Deep thinking (Homework)

- How can we solve the problem?

```
if ( noMilk ) {          if ( noMilk ) {
    if (noNote) {         if (noNote) {
        leave note;       leave note;
        buy milk;         buy milk;
        remove note;      remove note;
    }                     }
}                          }
```

27

## Often times, we have to wait for shared resources

28

## Producer & Consumer Problem (1-pool version)

- **Producer**: creates copies of a resource
- **Consumer**: uses up (destroys) copies of a resource. (may produce something else)
- **Buffer**: used to hold resource produced by producer before consumed by consumer
- **Synchronization**: keeping producer & consumer in sync
- Happens inside OS all the time (e.g. I/Os)
  - How about in real life?

29

## Producer & Consumer – solution using locks?

Producer

```
while (1) {  
  
    produce an item;  
  
    while (buffer is full);  
  
    insert item into buffer  
  
}
```

Consumer

```
While (1) {  
  
    while (buffer is empty);  
  
    remove an item;  
  
    consume the item  
  
}
```

30

## Often times, we have to wait for shared resources

- Busy waiting is a bad idea
- Checking resources itself needs to be in critical section
- Busy waiting inside CS even worse!
  - No one else can check!

→ Need a more powerful sync. primitive!

31

## Reading assignment

- Chapters 3, 5
- Chapter 6 (process synchronization)
  - Read up materials covered

32