

Semaphore Implementation

ECE595
Jan 27

Y. Charlie Hu



1

Synchronization Primitives provided by OS (language/compiler)



- Lock
 - Alone is not powerful enough
- Semaphore (incl. binary semaphore)
 - binary semaphore alone not enough
- Lock and condition variable
- Monitor (hide lock, still use condition variables)

4

Semaphore



```
wait(S) {                signal(S) {  
    while (S<=0);        S++;  
    S--;  
}
```

- Semantics:
 - S “counts” the number of “resources”
 - wait(S) signifies “consuming” one if there is any, otherwise wait!
 - signal(S) signifies “producing” one

6

Semaphore implementation

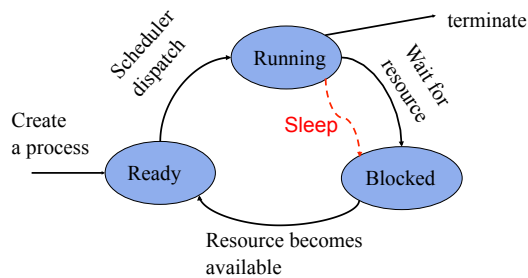


```
wait(S) {                signal(S) {  
    while (S<=0);        S++;  
    S--;  
}
```

- Can they be implemented in the user space?
 - An intuitive argument?
 - No existing hardware implements them directly
 - Scheduling/queuing cannot be easily done in HW
- Semaphore must be done in OS, typically with low-level synchronization support from hardware

8

[lec 2] Process State Transition



9

Uniprocessor solution

```

typedef struct {
    int count;
    queue q;
} semaphore;

void wait(semaphore s)
{
    if (s->count > 0) {
        s->count --;
        return;
    }
    add(s->q, current_process);
    sleep();
}

void signal(semaphore s)
{
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        process = removeFirst(s->q);
        wakeup(process);
        /* put process on Ready Q */
    }
}

/* Wake up a sleeping process == no op to counter
    
```

10

Challenge

- Need to make primitives atomic!
- What Mutex support do we have from HW on uniprocessor?
 - On uniprocessor, reads and writes are atomic

11

Uniprocessor solution

```

typedef struct {
    int count;
    queue q;
} semaphore;

void wait(semaphore s)
{
    if (s->count > 0) {
        s->count --;
        return;
    }
    add(s->q, current_process);
    /* imply sleep(); */
}

void signal(semaphore s)
{
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        process = removeFirst(s->q);
        wakeup(process);
        /* put process on ReadyQ */
    }
}
    
```

12

Uniprocessor solution

- What can cause the few lines to be not atomic?
- What causes context switches?
- Recall -- only way the OS dispatcher regains control is via **interrupts** (incl. HW interrupt like timer, SW interrupt like syscalls)
 - E.g. typing -> keyboard interrupt -> handler -> kernel -> user process

13


Uniprocessor solution: disable interrupts!

```
void wait(semaphore s)
{
    disable interrupts;
    if (s->count > 0) {
        s->count --;
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    enable interrupts;
    /* implying sleep(); */
}

void signal(semaphore s)
{
    disable interrupts;
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        process = removeFirst(s->q);
        wakeup(process);
        /* put process on Ready Q */
    }
    enable interrupts;
}
```

14

Generic shared-memory multiprocessor



15

What about multiprocessors?

- True concurrency – simultaneity!
 - Cache coherence in HW (fairly complicated)
 - For simplicity: a read sees most recent write
- Is turning off interrupt locally enough?
- Turning off interrupt on all other processors?
- Use atomic read/write and busy waiting, as in “too much milk”?

16

[Ice3] Yet Another Possible Solution?

```

process A
    leave noteA
    while (noteB)
        do_nothing;
    if (noMilk)
        buy milk;
    remove noteA

process B
    leave noteB
    if (noNoteA) {
        if (noMilk) {
            buy milk
        }
    }
    remove noteB
  
```

- Safe to buy
- If the other buys, I won't
- 2 things we dislike this solution?

17

A Multiprocessor solution?

```

void wait(semaphore s)
{
    disable interrupts;
    while (s->count == 0);

    s->count --;
    enable interrupts;
}

void signal(semaphore s)
{
    disable interrupts;
    s->count ++;
    enable interrupts;
}
  
```

- Does this work?

18

What about multiprocessors? (cont)

- Cannot just turn off interrupts
 - Turn off interrupt on local processor?
 - Turn off all other processors?
- Use atomic read/write and busy waiting, as in "too much milk"?
 - Concurrent read can happen
- Need more help from HW!
- What is the right, minimal HW support?
 - Hot research topic for a long time

20

[Lec2] A Possible Solution?

```

if ( noMilk ) {
    if (noNote) {
        leave note;
        buy milk;
        remove note;
    }
}

if ( noMilk ) {
    if (noNote) {
        leave note;
        buy milk;
        remove note;
    }
}
  
```

- process can get context switched after checking milk and note, but before leaving note
- Why does it work for human?

21

Read-modify-write on CISC

- Most CISC machines provide atomic *read-modify-write* instruction
 - read existing value
 - store back a new value
 - Example: *test-and-set* by IBM and others

```
X = TAS(&lock, 1);

read lock value V;
if 0 set lock to 1 else noop;
return V;
```

22

Using test-and-set for mutual exclusion (too-much milk)

- Implement a critical section on multiprocessor
 - Prevents 2 processes doing 0-to-1 transition simultaneously

```
global int lock = 0;
...
??????
...
critical section
...
??????
```

23

Use TAS to implement semaphores on multiprocessor

- For each semaphore, keep an extra integer (lock)

```
typedef struct {
    int lock; /* initially 0 */
    int count;
    queue q; /* queue of procs waiting on this semaphore */
} semaphore;
```

24

Use TAS to implement semaphores on multiprocessor?

```
void wait(semaphore s)
{
    disable interrupts;
    ???
    if (s->count > 0) {
        s->count --;
        ???
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    ???
    /* implying sleep(); */
    enable interrupts;
}

void signal(semaphore s)
{
    disable interrupts;
    ???
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        thread = removeFirst(s->q);
        wakeup(thread);
        /* put process on Ready Q */
    }
    ???
    enable interrupts;
}
```

25


Really Deep Thinking

- Why is this busy waiting not a concern?

26

Weaker mechanism on RISC [MIPS R4000 series]

- Load-linked** instruction: **LDL Rx,y**
 - loads Rx with a word from mem addr y
 - holds y in per-processor lock register
- Store operation to addr y (by any processor) resets all other processors' lock registers if containing addr y
- Store-conditionally** instruction: **STC Rx, y**
 - stores a word iff y matches the processor's lock register
 - indicates success (1) or failure (0)



27

Weaker mechanism on RISC [MIPS R4000 series]

- Load-linked** instruction: **LDL Rx,y**
 - loads Rx with a word from mem addr y
 - holds y in per-processor lock register
- Store operation to addr y (by any processor) resets all other processors' lock register if containing addr y
- Store-conditionally** instruction: **STC Rx, y**
 - stores a word iff y matches the processor's lock register
 - indicates success (1) or failure (0)

```
int lock=0;
...
while (ldl(&lock) == 1) || stc(&lock, 1) == 0);

...
critical section
...
lock = 0;
```

28

Use ldl/stc to implement semaphores on multiprocessor

```
void wait(semaphore s)
{
    disable interrupts;
    ???
    if (s->count > 0) {
        s->count --;
        ???
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    ???
    /* implying sleep(); */
    enable interrupts;
}

void signal(semaphore s)
{
    disable interrupts;
    ???
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        thread = removeFirst(s->q);
        wakeup(thread);
        /* put process on Ready Q */
    }
    ???
    enable interrupts;
}
```

29

Implementing Locks and Condition Variables



- Use the same mechanisms for
 - achieving atomicity and
 - avoiding busy waiting as in semaphore implementation

30

Reading assignment



- Chapter 6

32