

More Synchronization Problems, Wait-Free Synchronization

ECE595

Jan 30

Y. Charlie Hu



1

Roadmap



- Interprocess communication *with shared data*
 - Synchronization with locks, semaphores, condition var
 - Classic sync. problem 1: producer-consumer
 - Semaphore implementations (uniprocessor, multiprocessor)

Today:

- Classic sync. problems 2 & 3
- Wait-free synchronization

2

Classic Synchronization Problems



1. Producer-consumer problem (bounded buffer problem)
2. Readers-writers problem
3. Dining philosophers problem

4

Dining philosophers problem



Abstraction of concurrency-control problems

The need to allocate several resources among several processes while being *deadlock-free* and *starvation-free*



Classic Synchronization Problems

1. Producer-consumer problem (bounded buffer problem)
2. Readers-writers problem
3. Dining philosophers problem

6

Readers-Writers problem

Abstraction of concurrent access to shared data problem

- A data object is shared among multiple processes

Reader:

While (1) {

acq(mutex)

read();

rel(mutex)

}

Writer:

While (1){

acq(mutex)

write();

rel(mutex) /* abstraction */

}

7

Readers-Writers problem

Abstraction of concurrent access problem

- A data object is shared among multiple processes
- Allow concurrent reads, but exclusive writes
 - Implication: need to move read() and write() outside Critical Sec
 - Can we do it using lock+ flags?
 - Can we use semaphore to count readers/writers?

Reader:

acq(mutex)

????

rel(mutex)

read();

acq(mutex)

???

rel(mutex)

Writer:

acq(mutex)

????

rel(mutex)

write();

acq(mutex)

???

rel(mutex)

8

Readers-Writers problem

Abstraction of concurrent access problem

- A data object is shared among multiple processes
- Allow concurrent reads, but exclusive writes
- Solution needs lock, semaphores, and counting!
- Constraints:
 - Writers can only proceed if there are no active readers/writers
→ use semaphore OKtoWrite
 - Readers can proceed only if there are no active/waiting writers
→ use semaphore OKtoRead
 - To keep track of how many are reading / writing / waiting
→ use some shared variables, called state variables
 - Only one process manipulates state variable at once
→ use a lock Mutex

9

Readers-Writers problem (cont)

- State variables:
 - AR = number of active readers
 - WR = number of waiting readers
 - AW = number of active writers
 - WW = number of waiting writers
 AW is always 0 or 1
 AR and AW can not both be non-zero
- Initialization:
 - OKtoRead = 0;
 - OKtoWrite = 0;
 - Mutex = 1;
 - AR = WR = AW = WW = 0;
- Scheduling: writers get preference

10

Readers-Writers problem (cont)

- Reader

acq(mutex)

acq(mutex);

rel(mutex);

rel(mutex)
wait(OKtoRead);
read necessary data;

11

Readers-Writers problem (cont)

- Writer

acq(mutex)

acq(mutex);

rel(mutex);

rel(mutex)
wait(OKtoWrite);
write necessary data;

12

Readers-Writers problem (cont)

- Reader

acq(mutex)
if (???) {
 signal(OKtoRead);
 AR ++;
} else {
 WR ++;
}
rel(mutex)
wait(OKtoRead);
read necessary data;

acq(mutex);
AR--;
if (???) {
 V(OKtoWrite);
 AW ++;
 WW --;
}
rel(mutex);

13

Readers-Writers problem (cont)

- Writer

```
acq(mutex)
if (????) {
    signal(OKtoWrite);
    AW ++;
} else {
    WW ++;
}
rel(mutex);
wait(OKtoWrite);
write necessary data;
```

```
acq(mutex);
AW --;
if (????) {
    signal(OKtoWrite);
    AW ++;
    WW --;
} else while (????) {
    signal(OKtoRead);
    AR ++;
    WR --;
}
rel(mutex);
```

15

What happens if

- Reader enters and leaves system
- Write enters and leaves system
- Two writers enter system
- Two readers (a,b) enter system
- Writer(c) enters system and waits
- Reader(d) enters system and waits
- Readers(a,b) leave system, write(c) continues
- Write(c) leaves system, last reader(d) continues and leaves

17

Questions:

- In case of conflict between readers and writers, who gets priority?
 - Readers can get locked out
- Is the WW necessary in the writer's first if?
 - No: if there is a waiting writer, there must be an active writer or at least one active reader
- Can OKtoRead ever get greater than 1? What about OKtoWrite?
 - Yes, no
- Is the first writer to execute acq(mutex) guaranteed to be the first writer to access the data?
 - No, waiting writers can get granted in any order

18

Roadmap

- Interprocess communication *with shared data*
 - Synchronization with locks, semaphores, condition var
 - Classic sync. problem 1: producer-consumer
 - Semaphore implementations (uniprocessor, multiprocessor)

Today:

- Classic sync. problems 2 & 3
- Wait-free synchronization

19

[lec5] Uniprocessor solution: disable interrupts!

```

void wait(semaphore s)
{
    disable interrupts;
    if (s->count > 0) {
        s->count --;
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    enable interrupts;
    /* implying re-dispatch */
}

void signal(semaphore s)
{
    disable interrupts;
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        process = removeFirst(s->q);
        wakeup(process);
        /* put process on Ready Q */
    }
    enable interrupts;
}

```

20

[week3] Use TAS to implement semaphores on multiprocessor?

```

void wait(semaphore s)
{
    disable interrupts;
    while (tsa(&lock,1)==1);
    if (s->count > 0) {
        s->count --;
        lock=0;
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    lock=0;
    sleep(); /* re-dispatch */
    enable interrupts;
}

void signal(semaphore s)
{
    disable interrupts;
    ???
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        thread = removeFirst(s->q);
        wakeup(thread);
        /* put process on Ready Q */
    }
    ???
    enable interrupts;
}

```

Do we still need to disable interrupts?

21

Wait-free Synchronization

- Finally we need tsa or ldl^stc anyway to implement sync. primitives (on multiprocessors)
- Can we design data structures in a way that allows safe concurrent accesses?
 - no OS-supplied mutual exclusion necessary
 - no possibility of deadlock
 - only using tsa / ldl^stc
 - no busy waiting


22

Simple example – Queue insertion

```

typedef struct {
    QItem *item;
    QElem *next;
} QElem;


```



23

Simple example – Queue insertion

```
typedef struct {
    QItem *item;
    QElem *next;
} QElem;
```



24

Singly-linked Queue Insertion

```
QElem *queue;

void Insert(item) {
    QElem *new = malloc(sizeof(QElem));
    new->item = item;
    new->next = queue;
    queue = new;
}
```

25

Wait-free Synchronization

- Design data structures in a way that allows safe concurrent accesses
 - no mutual exclusion (lock acquire & release) necessary
 - no possibility of deadlock
- Approach: use a single atomic operation to
 - commit all changes
 - move the shared data structure from one consistent state to another

26

Read-modify-write on CISC

- Most CISC machines provide atomic *read-modify-write* instruction
- Assume a test-and-set instruction

```
X = TAS(addr, old_value, new_value);

read value V at addr;
if (V == old_value) set it to new_value;
return V;
```

27

Singly-linked Queue Insertion using TSA?



```
QElem *queue;

void Insert(item) {
    QElem *new = malloc(sizeof(QElem));
    new->item = item;
    new->next = queue;
    queue = new;
}
```

28

Limitation



- Example only works for simple data structures where changes can be committed with *one store instruction*
- What about more complex data structures?

30

More General Approach



- Maintain a pointer to the “master copy” of the data structure
- To modify,
 1. remember current value of the master pointer
 2. copy shared data structure to a scratch location
 3. modify copy
 4. *atomically*:
 - verify that master pointer has not changed
 - write pointer to refer to new master
 5. if verification fails (another process interfered), start over at step 1
- Downside?
 - When does it work reasonably well?

31

Reading



- Chapter 6

32