

Communication with Messages

ECE595

Feb 1

Y. Charlie Hu



1

Roadmap



- Interprocess communication *with shared data*
 - Synchronization with locks, semaphores, condition var
 - Classic sync. problem 1: producer-consumer
 - Semaphore implementations (uniprocessor, multiprocessor)
 - Classic sync. problems 2 & 3
 - Wait-free synchronization

Today:

- Interprocess communication *with messages*
- Project 2

2


Inter-process Communication with Messages



- Messages provide for communication *without shared data*
 - One process or the other owns the data, (guaranteed) never two at the same time
 - Think about usmail

3

Big Picture



4

Why use messages?

- Many types of applications fit into the model of processing a sequential flow of information
- Communication across address spaces – no side effects
 - Less error-prone
 - They might have been written by diff programmers who aren't familiar with code
 - They might not trust each other
 - They may not be running on different machines!
 - Examples?

5

Message Passing API

- Generic API
 - `send(mailbox, msg)`
 - `recv(mailbox, msg)`
- What is a mailbox?
 - A *buffer* where messages are stored between the time they are sent and the time when they are received
- What should “msg” be?
 - Fixed size msgs
 - Variable sized msgs: need to specify sizes

6

Buffering leads to design options

- When should `send()` return?
- When should `recv()` return?

7

Send

- *Fully Synchronous*
 - Will not return until data is received by the receiving process
- *Synchronous*
 - Will not return until data is received by the mailbox
 - Block on mailbox full
- *Asynchronous*
 - Return immediately
 - Completion
 - Require the application to check status (appl polls)
 - Notify the application (OS sends interrupt)
 - Block on mailbox full

8

Receive

- *Synchronous*
 - Return data if there is a message
 - Block on empty buffer
- *Asynchronous*
 - Return data if there is a message
 - Return status if there is no message (probe)

9

OS Implementation

- What is the conceptual problem for OS implementation here?
 - Assume sender and receiver are on the same machine


10

No Buffering

- Sender must wait until the receiver receives the message
- Rendezvous on each message

11

Mailbox - Bounded Buffer



- Buffer
 - Has fixed size
 - Is a FIFO
 - Variable size message
- Multiple producers
 - Put data into the buffer
- Multiple consumers
 - Remove data from the buffer
- Blocking operations
 - Sender waits if not enough space
 - Receiver waits if no message
- Synchronization
 - Using lock/condition variable (or semaphore)

12

Direct Communication

- Each process must name the sending or receiving process
- A communication link
 - is set up between the pair of processes
 - is associated with exactly two processes
 - exactly one link between each pair of processes

```
P: send( process Q, msg )
Q: recv( process P, msg )
```

13

Producer-Consumer Problem with Message Passing

```
Producer(){
  while (1) {
    ...
    produce item
    ...
    send( consumer, item);
  }
}
```

```
Consumer(){
  while (1) {
    recv( producer, item );
    ...
    consume item
    ...
  }
}
```

14

Indirect Communication

- Use a “mailbox” or “ports” to allow **many-to-many communication**
 - Mailbox typically owned by the OS
 - Requires open/close a mailbox before allowed to use it
- A “link”
 - is set up among processes only if they have a shared mailbox
 - Can be associated with more than two processes

```
P: open (mailbox); send( mailbox, msg);
   close(mailbox)
Q: open (mailbox); recv( mailbox, msg );
   close(mailbox)
```

15

The big debate in parallel computing: Messaging vs. Sharing Data

- Two programming models are equally powerful
- But result in very different-looking programming styles
- Most people find shared-data programming easier to work with
 - Debugging?
- What about machines that do not share memory?
 - Can be simulated in software [SDSM – hot topic in 80-90’ s]
 - But often not as efficient as message passing

16