

Sharing Main Memory, Segmentation


ECE595

Feb 8

Y. Charlie Hu

1

Connecting the dots



4

The big picture

- a.out needs address space for
 - text seg, data seg, and (hypothetical) heap, stack
- A running process needs phy. memory for
 - text seg, data seg, heap, stack
- But no way of knowing where in phy mem at
 - Programming time, compile time, linking time
- **Best way out?**
 - Make agreement to divide responsibility
 - Assume address starts at 0 at prog/compile/link time
 - OS needs to work hard at loading/runing time


5

Big picture (cont)

- OS deals with physical memory
 - Loading
 - Sharing physical memory between processes
 - Dynamic memory allocation

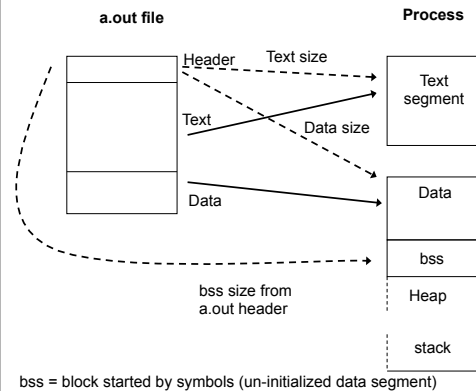
6

Connecting the dots



7

Loading




8

Dynamic memory allocation during program execution

- Stack: for procedure calls
- Heap: for malloc()
- Both dynamically growing/shrinking
- Assumption for now:
 - Heap and stack are fixed size
 - OS has to worry about loading 4 segments per process:
 - Text
 - Data
 - Heap
 - stack

9

1. Simple uniprogramming: Single segment (code, data, stack heap) per process



10

Simple uniprogramming: Single segment per process



- Highest memory holds OS
- Process is allocated memory starting at 0, up to the OS area
- When loading a process, just bring it in at 0
 - virtual address == physical address!
- Examples:
 - early batch monitor which ran only one job at a time
 - if the job wrecks the OS, reboot OS
 - 1st generation PCs operated in a similar fashion
- Pros / Cons?

11

Multiprogramming



- Want to let several processes coexist in main memory

12

Issues in sharing main memory



- **Transparency:**
 - Processes should not know memory is shared
 - Run regardless of the number/locations of processes
- **Safety:**
 - Processes mustn't be able to corrupt each other
- **Efficiency:**
 - Both CPU and memory utilization shouldn't be degraded badly by sharing

13

2. Simple multiprogramming




With **static software memory relocation**, no protection, 1 segment per process:

- Highest memory holds OS
- Processes allocated memory starting at 0, up to the OS area
- When a process is loaded, **relocate** it so that it can run in its allocated memory area
 - How? (use symbol table and relocation info)
- Analogy to linking?


14

Simple multiprogramming: Single segment per process, static relocation



15

Simple multiprogramming: Single segment per process, static relocation



16

Simple multiprogramming: Single segment per process, static relocation

- 4 drawbacks
 1. No protection
 2. Low utilization -- Cannot relocate *dynamically*
 - Binary is fixed (after loading)
 - Cannot do anything about holes
 3. No sharing -- Single segment per process
 - Cannot share part of process address space (e.g. text)
 4. Entire address space needs to fit in mem
 - Need to swap whole, very expensive!

17

What else can we do?

- Already tried
 - Compile time / linking time
 - Loading time
- Let us try execution time!

18


3. Dynamic memory relocation

- Instead of changing the address of a program before it's loaded, change the address dynamically *during every reference*
 - Under dynamic relocation, each program-generated address (called a *logical address* or *virtual address*) is translated in hardware to a *physical* or *real address*

Can this be done in software?

19


Translation overview



- Actual translation is in hardware (MMU)
- Controlled in software
- CPU view
 - what program sees, virtual addresses
- Memory view
 - physical memory addresses

20


3.1 Base and bound



- Built in Cray-1 (1976)
- A program can only access physical memory in [base, base+bound]
- On a context switch: save/restore base, bound registers
- Pros:
 - simple, fast, cheap
 - Can relocate segment

22


3.1 Base and bound



- The essence:
 - A level of indirection
 - $\text{Phy. Addr} = \text{Vir. Addr} + \text{base}$
- How to relocate segment in physical memory?
 - From Base 1 to Base 2?

23

Base and bound



- Cons:

- Only one segment
- How can two processes share code while keeping private data areas (shared editors)?

24

What have we achieved?

- 4 drawbacks

1. No protection
2. Low utilization -- Cannot relocate dynamically
 - Cannot do anything about holes
3. No sharing -- Single segment per process
 - Cannot share part of process address space (e.g. text)
4. Entire address space needs to fit in mem
 - Need to swap whole, very expensive!

25

Readings

- Chapter 8

26