

Sharing Main Memory, Segmentation (cont)

ECE595

Feb 10

Y. Charlie Hu



1


Outline

- Segmentation
- External fragmentation & Generic dynamic allocation problem
- OS implementation of segmentation



2

Connecting the dots



3

3. Dynamic memory relocation


- Instead of changing the address of a program before it's loaded, change the address dynamically *during every reference*
 - Under dynamic relocation, each program-generated address (called a *logical address* or *virtual address*) is translated in hardware to a *physical* or *real address*

Can this be done in software?



4


Translation overview



- Actual translation is in hardware (MMU)
- Controlled in software
- CPU view
 - what program sees, virtual addresses
- Memory view
 - physical memory addresses

5

3.1 Base and bound



- Built in Cray-1 (1976)
- A program can only access physical memory in [base, base+bound]
- On a context switch: save/restore base, bound registers
- Pros:
 - simple, fast, cheap
 - Can relocate segment


6

All problems in computer science can be solved by another level of indirection.

David Wheeler

7


3.1 Base and bound



- The essence:
 - A level of indirection
 - $\text{Phy. Addr} = \text{Vir. Addr} + \text{base}$
- How to relocate segment in physical memory?
 - From Base 1 to Base 2?

8

3.1 Base and bound



- Cons:
 - Only one segment
 - How can two processes share code while keeping private data areas (shared editors)?


9

What have we achieved?

- 4 drawbacks
 1. No protection
 2. Low utilization -- Cannot relocate dynamically
 - Cannot do anything about holes
 3. No sharing -- Single segment per process
 - Cannot share part of process address space (e.g. text)
 4. Entire address space needs to fit in mem
 - Need to swap whole, very expensive!

10

3.2 Multiple Segments



- Have a table of (seg, size)
- Further protection: each entry has (nil, read, write, exec)
- On a context switch: save/restore the table (or a pointer to the table) in kernel memory

11

How does this allow 2 processes to share code segment?

12

Segmentation example

text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

0320: bar:

01C2: jmp 01F4

01E0: call 0320

Data segment [0x1000, 0x16A0]

01F4: X:

15DC: _Y:

2-bit segment number, 12-bit offset

Segment Base Bounds RW

0	4000	4B0	10
1	0	6A0	11
2	3000	FFF	11
3	--	--	00

13

Segmentation example

text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

0320: bar:

01C2: jmp 01F4

01E0: call 0320

Data segment [0x1000, 0x16A0]

01F4: X:

15DC: _Y:

2-bit segment number, 12-bit offset

Segment Base Bounds RW

0	4000	4B0	10
1	0	6A0	11
2	3000	FFF	11
3	--	--	00

14

→ Where is 01F4 in physical memory?

Segmentation example

text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

0320: bar:

01C2: jmp 01F4

01E0: call 0320

Data segment [0x1000, 0x16A0]

01F4: X:

15DC: _Y:

2-bit segment number, 12-bit offset

Segment Base Bounds RW

0	4000	4B0	10
1	0	6A0	11
2	3000	FFF	11
3	--	--	00

15

→ Where is 15DC in physical memory?

Segmentation example

text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

0320: bar:

01C2: jmp 01F4

01E0: call 0320

Data segment [0x1000, 0x16A0]

01F4: X:

15DC: _Y:

2-bit segment number, 12-bit offset

Segment Base Bounds RW

0	4000	4B0	10
1	0	6A0	11
2	3000	FFF	11
3	--	--	00

16

→ Suppose SP is initially 265C. Where is it in physical mem?

Segmentation example

text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

0320: bar:

01C2: jmp 01F4

01E0: call 0320

Data segment [0x1000, 0x16A0]

01F4: X:

15DC: _Y:

2-bit segment number, 12-bit offset

Segment Base Bounds RW

0 4000 4B0 10

1 0 6A0 11

2 3000 FFF 11

3 -- -- 00

→ which portions of the virtual and physical address spaces are used by this process?

17

Pros/cons of segmentation

• Pros:

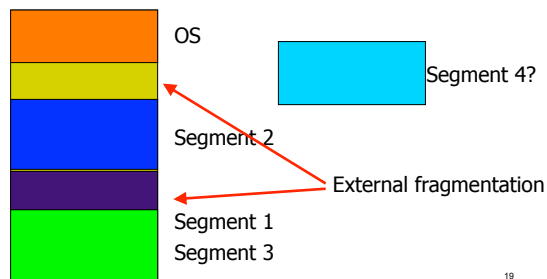
- Process can be split among several segments
 - Allows sharing
- Segments can be assigned, or swapped independently

• Cons:

- **External fragmentation:** many holes in physical memory
 - Also happens in base and bound scheme

18

Simple multiprogramming: Single segment per process, static relocation



19

What fundamentally causes external fragmentation?

1. Segments of many different sizes
2. Each has to be allocated contiguously

20

Can we solve **Dynamic** memory allocation problem?



- Problem: External fragmentation caused by holes too small
- How much can a smart allocator help?
 - The allocator maintains a free list of holes
 - Allocation algorithms differ in how to allocate from the free list

21

Dynamic allocation algorithms



- Best fit: allocate the smallest chunk big enough
- First fit: allocate the first chunk big enough
 - Rotating first fit
- Is best fit necessarily better than first fit?
 - Example: 2 free blocks of size 20 and 15
 - If allocation ops are 10 then 20, which one wins?
 - If ops are 8, 12, then 12, which one wins?

22

Dynamic allocation algorithms



- Analysis shows
 - First fit tends to leave average-size holes
 - Best fit tends to leave some very large holes, very small holes
- Knuth claims that if storage is close to running out, it will run out regardless of which scheme is used
 - Pick the easiest or most efficient (e.g. first fit)

23

Segmentation: OS implementation



- Keep segment table in PCB
- When creating process, allocate space for segments, fill in PCB bases/bounds
- When process dies, return physical space used by segments to free pool
- Context switch?
 - Saves old segment table / Loads new segment table to MMU
 - What about context switch of threads?
 - True-or-false: CS between threads of same process cheaper than CS between processes

[lec2] Kernel data structure: Process Control Block (Process Table)

- Process management info
 - State (ready, running, blocked)
 - PC & Registers, parents, etc
 - CPU scheduling info (priorities, etc.)
- Memory management info
 - Segments, page table, stats, etc
- I/O and file management
 - Communication ports, directories, file descriptors, etc.



Managing segments (cont)

To enlarge a segment:

- See if space above segment is free. If so, just update the bound and use that space
- Or, move this segment to disk and bring it back into a larger hole (or maybe just copy it to a large hole)

26

Managing segments (cont)

- When there is no space to allocate a new segment:
 - Compact memory – how?

27

Summary: Evolution of Memory Management (so far)

Scheme	How	Pros	Cons
Simple uniprogramming	1 segment loaded to starting address 0	Simple	1 process 1 segment No protection
Simple multiprogramming	1 segment relocated at loading time	Simple, Multiple processes	1 segment/process No protection External frag.
Base & Bound	Dynamic mem relocation at runtime	Simple hardware, Multiple processes Protection	1 segment/process, External frag.
Multiple segments	Dynamic mem relocation at runtime	More hardware, Protection, multi segs/process	External frag.

28

