

Paging

ECE595

Feb 13

Y. Charlie Hu


2

Review: sharing main memory

- Linker generated a.out, assuming address starting from 0
- OS needs to load it into physical memory
 - Code
 - Data
 - Stack
 - Heap

9


[lec13] The big picture



10

Review: sharing main memory


- Simple uniprogramming



11

Review: sharing main memory

- Simple multiprogramming



4 drawbacks?

12

Review: sharing main memory


- Simple multiprogramming – 4 drawbacks
 - No protection
 - Low utilization -- Cannot relocate dynamically
 - Cannot do anything about holes
 - No sharing -- Single segment per process
 - Entire address space needs to fit in mem
 - Need to swap whole, very expensive!

13

Fix drawback #1 & #2: Dynamic memory relocation

- Instead of changing the address of a program when it's loaded, change the address dynamically *during every reference*
 - Under dynamic relocation, each program-generated address (called a *logical address* or *virtual address*) is translated in hardware to a *physical* or *real address*


Dynamic memory relocation



- Actual translation is in hardware (MMU)
 - why?
- Controlled in software
- CPU view
 - what program sees, virtual addresses
- Memory view
 - physical memory


15

[lec13] The big picture



16

Fix drawback # 1 & #2 – base and bound



Indirection!

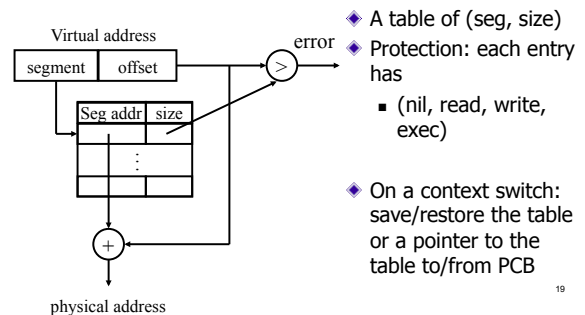
17

Review: sharing main memory

- Simple multiprogramming – 4 drawbacks
 1. No protection
 2. Low utilization -- Cannot relocate dynamically
 - Can relocate now; but is frequent relocation desirable?
 3. **No sharing -- Single segment per process**
 4. Entire address space needs to fit in mem
 1. Need to swap whole, very expensive!

18

Fix drawback #3 & #4 – allow multiple segments



19

Pros/cons of segmentation

- Pros:
 - Process can be split among several segments
 - Allows sharing (how?)
 - Segments can be allocated/swapped independently
 - Still allocate/swap each segment as a whole
- Cons:
 - **External fragmentation**: many holes in physical memory
 - Also happens in base and bound schemes
 - Can relocate, but is it desirable / easy?


20

What fundamentally causes external fragmentation?

- Segments of many different sizes
 - Each has to be allocated contiguously
-
- “Million-dollar” question:
Physical memory is precious.
Can we limit the waste to a single hole of X bytes?

21

Virtual pages / physical pages




22

Paging

- Goal:
 - to make allocation and swapping easier (time)
 - to reduce memory fragmentation (space)
- Key idea:
 - Make all chunks of memory the same size, called **pages**
- Implementation:
 - For each process, a **page table** defines the base address of each of that process' pages along with existence and read/write bits
 - Translation?

23

Paging




- Context switch
 - similar to the segmentation scheme
- Pros:
 - easy allocation, keep a free list
 - easy to swap
 - easy to share

24

Deek thinking: Paging implementation

- Translation: table lookup and bit substitution
- Why is this possible?
- Why cannot we do the same in segmentation?




25

How many PTEs do we need? (assume page size is 4096 bytes)

- Worst case for 32-bit address machine?
- What about 64-bit address machine?

26

Paging implementation – how does it really work?



- Where to store page table?
- How to use MMU?
 - PT too large to fit into MMU
 - PT kept in mem and MMU only has base addresses
 - What does MMU have to do?
- Page size?
 - Small page -> big table
 - 32-bit with 4k pages
 - Large page -> small table but large internal fragmentation
 - e.g., data seg is 9 Kbytes

27

Paging vs. segmentation

- Segmentation:
 - External fragmentation
 - Complicated allocation, swapping
 - + Small segmentation table
- Paging
 - Internal fragmentation
 - + Easy allocation, swapping
 - Large page table

28

Deep thinking

- Why does the page table have to be contiguous in the physical memory?
 - Why did a segment have to be contiguous in memory?
- For a 4GB virtual address space, we just need 1M PTE (~4MB), what is the big deal?
- My PC has 512MB, why do we need PTEs for the entire 4GB address space?

29


Page table

- The page table has to be consecutive in mem
 - Potentially large
 - Consecutive pages in mem hard to find
- How can we be flexible?

“All computer science problems can be solved with an extra level of indirection.”

30

Two-level page tables



What does this buy us?

31

Multi-level page tables

- 3 Advantages?

The power of an extra level of indirection!

32



Reading assignment

- Chapter 8

33

