

File Caching

ECE595

Mar 24

Y. Charlie Hu



[week11] UFS concurrency semantics



- What happens when two processes try to write to the same file?
 - What is the programmer's intent?
- What needs to be made atomic in FS impl?
 - AllocateBlock(); FreeBlock()
 - Write() – AllocateBlock() and update inode
- What are naturally atomic?
 - WriteRawData(); ReadRawData()
- Analogy?

2

[week9] Disk Allocation revisited – many low level details



- How to keep blocks for a file together?
- How about inode and data blocks for a file?
 - It is a good idea to keep them close?
 - If so, how?
- How about files in the same directory?
 - e.g. make

3

True or False



- On Unix, a kernel process can write a dir just like reading/writing an ordinary file
- On Unix, a user process can write a dir just like reading/writing an ordinary file, assuming the user has the read/write permission

4

Roadmap

- Functionality (API)
 - Basic functionality
 - Disk layout
 - File operations (open, read, write, close)
 - Directories
- Performance
 - Disk allocation
 - Buffer cache
 - File system interface
 - Disk scheduling
- Reliability
 - FS level
 - Disk level: RAID

5

“Principle of locality” once more

- Locality of reference in file accesses
 - Yet another application of the principle of locality
 - What were the earlier instances in this class?
- Keep a number of disk blocks in “the much faster” memory
 - when accessing disk, check the cache first!
- File system buffer caches are maintained in software

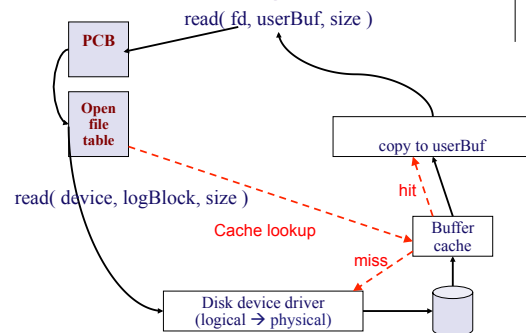
6

How many disk blocks to cache?

- Fixed portion of main memory (BSD)
- Variable portion of main memory (modern Unix) -- processes and file system compete for physical memory
- Pros/cons?

7

[week8] Reading A Block




Modern disk drives are addressed as large one-dimensional arrays of logical blocks

8

Read operations in presence of buffer cache

- `read(fd, buf, n)`
 - On a hit
 - copy from the buffer cache to a user buffer
 - On a miss
 - replacement if necessary
 - read a file block into the buffer cache




- Concurrent reads?

9

Write operations: Maintaining Consistency

- `write(fd, buffer, n)`
 - On a hit
 - write to buffer cache
 - On a miss?



- Concurrent writes?

10

File persistence under file caching

- Problem: fast cache memory is **volatile**, but users expect disk files to be **persistent**
 - In the event of a system crash, dirty blocks in the buffer cache are lost !
- Example 1: creating `"/dir/a"`
 - Allocate inode (from free inode list) for "a"
 - Update parent dir content – add `<"a", inode#>` to "dir"

11

File persistence under file caching

- Problem: fast cache memory is volatile, but users expect disk files to be persistent
 - In the event of a system crash, dirty blocks in the buffer cache are lost !
- Example 2: append a block to a file
 - Allocate data block (from free block list)
 - Update inode content
 - Write to new data block

12

File persistence under file caching



- Solution 1: use **write-through** cache
 - Modifications are written to disk immediately
 - (minimize "window of opportunities")
 - No performance advantage for disk writes
- Example 2: append a block to a file
 - Allocate data block (from free block list)
 - Update inode content
 - Write to new data block

13

File persistence under file caching



- Solution 2: limit potential data loss (Unix)
 - **Write-through** caching for metadata (inodes, directories, free block list)
 - **Write back**
 - dirty data blocks after no more than 30 seconds
 - all dirty blocks during file close
 - Worse case damage?
- Example 2: append a block to a file
 - Allocate data block (from free block list)
 - Update inode content
 - Write to new data block

14

File caching implementation



- Two major issues
 - Buffer cache replacement
 - What problem does this resemble?
 - How are they different?
 - Implications?
 - Competition with VM for main memory
 - Static partitioning during kernel configuration (BSD)
 - Dynamic adjustment of partitioning during runtime,
 - e.g. keep miss frequencies of VM and buffer cache, and try to balance them (e.g. Linux)

15

Buffer cache replacement



- A classic OS research topic
- New ideas still come out
- Recency / frequency based
 - (exact) LRU, MRU, LRU-K, FBR, LRFU, etc.
- Pattern based - manual
 - User inserted / compiler generated
- Pattern based - automatic
 - DEAR: per appl pattern classification
 - UBM: per file pattern classification
 - PCC: per call-site classification [Gniady/Butt/Hu OSDI '04]

16

Other performance optimizations



- Read-ahead (e.g. Linux)
 - For sequential access, read the requested block and the following N blocks together (why is this a good idea?)
- Write-behind:
 - Start disk write, but don't make application wait until the disk operation completes
- Allow overlap of a process's computation with its own disk I/O (e.g. AIO in FreeBSD)