

## File System Interface

ECE595

Mar 29

Y. Charlie Hu



## [week8] File System API



- OS provides the file system abstraction
- How do application processes access the file system?

2

## Can we directly write() to a dir file?



- ... using system calls (e.g. write)
- This is a permission question, not a privilege question
- The answer is no
- Explanation:
  - Any process independent of file system shouldn't be allowed to directly write in a directory the way they can write in files
  - Since such processes don't know how exactly the internals of directory are laid out. They should only see files, create files, and delete files.

3

## System Calls to UNIX File Systems



- 19 system calls into 6 categories:

Return file desp.	Assign inodes	Set file attr.	Process input/output	Change file system	Modify view of file system
open close creat pipe dup	creat link unlink	chown chmod stat fstat	read write lseek	mount umount	chdir chroot

4

## Roadmap

- Functionality (API)
  - Basic functionality
    - Disk layout
    - File operations (open, read, write, close)
  - Directories
- Performance
  - Disk allocation
  - Buffer cache
  - File System interface
  - Disk scheduling
- Reliability
  - FS level
  - Disk level: RAID

5

## [week1] Roadmap for ECE469 lectures

Introduction to OS components

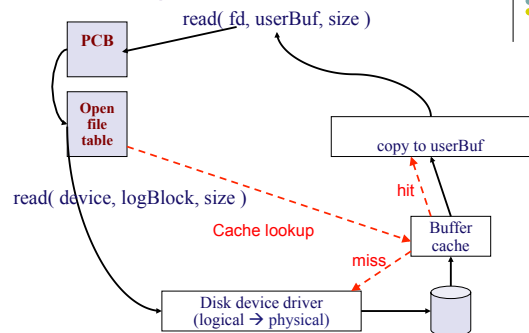
Individual components

- Process management
- Memory management
- File management
- Secondary-storage management
- (Device management)
- (shell)
- (Networking)

Hardware support for OS interspersed (before relevant topics)

6

## Reading A Block



Modern disk drives are addressed as large one-dimensional arrays of logical blocks

7

## File System Interface

- How do application programs typically access file data?
  - Explicit read/write operations (conventional)

8

## Read / Write Interface

- File data is explicitly copied between disk file and process memory
- Programs cannot directly access file data

```


FileDescriptor fhandle;
int offset, length;
char *buffer = (char *) malloc(1000000);

fhandle = open("pathname");
read(fhandle, offset, buffer, length);
{read file data in buffer to do important computation};
Write(fhandle, offset, buffer, length)
close(fhandle);
    
```

9


## Read / Write Interface – Problem 1

- Potential for **double copies** in mem



10


## Essence of Demand Paging



Disk is the backing store

11


## [week6] Demand Paging, Page Fault Handling



12


## Read / Write Interface – Problem 2

- Potential for **double paging** on disk
  - process pages containing file data are paged out to paging space, leading to redundant copies of file data on disk



13


## Essence of Demand Paging



Disk is the backing store

14

## What if



15

## Memory-mapped Files

- File is “mapped” into application’s address space
  - by initializing virtual memory so that the file (directly) serves as backing store for a region of the application's address space

16

## Memory-mapped Files (cont')

- Elegant integration of file system and virtual memory

```
FileDescriptor fhandle;  
int offset, length;  
char *address;  
  
fhandle = open("pathname");  
mmap(fhandle, offset, address, length);  
{read/write file data by accessing memory range  
  [address, address + length]};  
munmap(address, length);  
close(fhandle);
```

This is like after  
address = malloc()

17

## Implementation of mmap()

18

## Memory-mapped Files

- File is "mapped" into application's address space
  - by initializing virtual memory so that the file (directly) serves as backing store for a region of the application's address space
- File data is *demand paged* upon access to the mapped file
  - No double paging on disk
- Memory-mapped files do not go through buffer cache
  - No double copy in mem
    - Program accesses file data directly

19

## Effects and Semantics of Memory-mapped Files

- Processes that map the same file share **physical memory** that caches file data
- Writes may not be immediately written to the file on disk
  - Update periodically
  - Closing the file results in writing all to disk and removing the VM mapping

20

## Fun with memory-mapped files



- Inter-process communication
  - Virtual addresses of diff processes mapped to the same file
- File copying as Memory copying
  - Map files to virtual addresses
  - Do memory copying

21

## Reading



- Chapters 11-12

22