

Deadlocks

ECE595
Apr 12

Y. Charlie Hu



1

Review: Preemptive CPU Scheduling



- What is in it?
 - Mechanism + policy
 - Mechanisms fairly simple
 - Policy choices harder

2

Review: Evolution of CPU scheduling policies



- Don't know future → optimal policy is hard
- FIFO, Round-Robin, SJF all have merits
 - Tradeoffs are tricky to analyze
 - occasionally we can prove things
- Need a general framework to encompass all
 - Priority scheduling
- But coming up with priorities is tricky
 - Multiple queue scheduling
- But statically assigning queues not flexible
 - multi-level feedback queue scheduling

3

Quiz – one-sentence answer



- (Assume all jobs are doing CPU only, and all jobs arrived in a burst at time 0). Under what job arrival order, does FIFO give the worst avg. turnaround time?
- Can you prove it?


4

Quiz – True or False

- “A CPU scheduling algorithm that minimizes Avg. turnaround time cannot lead to starvation.”
- “Among all CPU scheduling algorithms, Round Robin always gives the worse average turnaround time.”

5

[week4] Dining Philosophers Problem (using semaphores)



Phil(i)

think;

wait(i.right)
wait(i.left);

eat;

signal(i.left);
signal(i.right);

How did we solve it?

6

Deadlock Example

- A law passed by the Kansas legislature early in the 20th century (in part):

“When two trains approach each other at a crossing, both come to a full stop and neither should start up again until the other has done.”

7

Deadlock Example



8

Deadlocks

- **Definition:** in a multiprogramming environment, a process is **waiting** forever **for** a resource held by another **waiting** process



- Topics:
 - Conditions for deadlocks
 - Strategies for handling deadlocks

11

System Model

- Resources
 - Resource types R_1, R_2, \dots, R_m
 - CPU cycles, memory space, I/O devices, mutex
 - Each resource type R_i has W_i instances
 - **Preemptable:** can be taken away by scheduler, e.g. CPU
 - **Non-preemptable:** cannot be taken away, to be released voluntarily, e.g., mutex, disk, files, ...
- Each process utilizes a resource as follows:
 - request
 - use
 - release

12


Resource-Allocation Graph

- A set of vertices V and a set of edges E
- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the **processes** in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all **resource types** in the system
- E is partitioned into two types
 - **request edge** – directed edge $P_i \rightarrow R_j$
 - **assignment edge** – directed edge $R_j \rightarrow P_i$

13

Resource-Allocation Graph (Cont.)


- Process
- Resource type with 4 instances
- P_i requests instance of R_j
- P_i is holding an instance of R_j



14


Example of a Resource Allocation Graph – one instance per type

- What happens if there is a cycle in the resource allocation graph?




15

Example of a Resource Allocation Graph – multiple instances / type




16

Resource Allocation Graph with a cycle – is there a deadlock?



17

Resource Allocation Graph with a cycle – is there a deadlock?



18

Basic Facts

- If graph contains no cycles \Rightarrow no deadlock
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, **possibility** of deadlock


19

4 Necessary Conditions for Deadlock

- **Mutual exclusion**
 - Each resource instance is assigned to exactly one process
- **Hold and wait**
 - Holding at least one and waiting to acquire more
- **No preemption**
 - Resources cannot be taken away
- **Circular chain of requests**

Program behavior

Resource nature




Eliminating **any** condition eliminates deadlock!

20

Eliminate Competition for Resources?

- If running A to completion and then running B, there will be no deadlock
- Generalize this idea for all processes?
- Is it a good idea?



21

Four Possible Strategies

1. Ignore the problem
 - It is user's fault
 - used by most operating systems, including UNIX
2. Detection and recovery (by OS)
 - Fix the problem after occurring
3. Dynamic avoidance (by OS, programmer help)
 - Careful allocation
4. Prevention (by programmer, practically)
 - Negate one of the four conditions

22

4 Necessary Conditions for Deadlock



Resource nature

- *Mutual exclusion*
 - Each resource instance is assigned to exactly one process
- *Hold and wait*
 - Holding at least one and waiting to acquire more
- *No preemption*
 - Resources cannot be taken away
- *Circular chain of requests*

Program behavior

Eliminating **any** condition eliminates deadlock!

23

4.1 Prevention: Remove Mutual Exclusion



- Some resources can be made sharable
 - Read-only files, memory, etc
- Some resources are not sharable
 - Printer, tape, mutex, etc
- Dining philosophers problem?

24

4.3 Prevention: Preemption (w/o changing app)



- Make scheduler aware of resource allocation
- Method
 - If a request from a process holding resources cannot be satisfied, preempt the process and release all resources
 - Schedule it only if the system satisfies all resources
- Applicability?
 - Preemptable resources:
 - CPU registers, physical memory
 - Difficult for OS to understand app intention
- Dining philosophers problem?

25

4 Necessary Conditions for Deadlock



Resource nature

- *Mutual exclusion*
 - Each resource instance is assigned to exactly one process
- *Hold and wait*
 - Holding at least one and waiting to acquire more
- *No preemption*
 - Resources cannot be taken away
- *Circular chain of requests*

Program behavior

Eliminating **any** condition eliminates deadlock!

26

4.2 Prevention: (change app) Remove Hold and Wait

- Two-phase locking
 - Phase I:
 - Try to lock all needed resources at the beginning
 - Phase II:
 - If successful, use the resources & release them
 - If not, release all resources and start over
- This is how telephone company prevents deadlocks
- Dining philosophers problem? (use TSA)
- 2 Problems with this approach?

27

[week3] Using test-and-set for mutual exclusion (too-much milk)


- Implement a critical section on multiprocessor
 - Prevents 2 processes doing 0-to-1 transition simultaneously

```
global int lock = 0;
...
while (TAS(&lock, 1) == 1);
...
critical section
...
Lock = 0;
```

28

4.4 Prevention: (change app) No Circular Wait

- Impose some order of requests for all resources
- How?
- Dining philosophers problem?
- Can we prove it always works?
- How is this different from two-phase locking?



29

Four Possible Strategies

1. Ignore the problem
 - It is user's fault
 - used by most operating systems, including UNIX
2. Detection and recovery (by OS)
 - Fix the problem afterwards
3. Dynamic avoidance (by OS & programmer)
 - Careful allocation
4. Prevention (by programmer & OS)
 - Negate one of the four conditions

30

Reading assignment

- Read chapter 7

